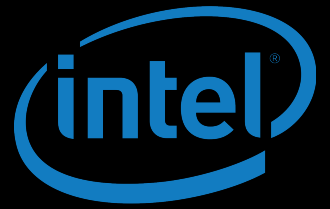

Reducing the Storage Overhead of Main-Memory OLTP Databases with Hybrid Indexes

SIGMOD'16

Huanchen Zhang

Co-authors:
David G. Andersen
Andrew Pavlo
Michael Kaminsky
Lin Ma
Rui Shen

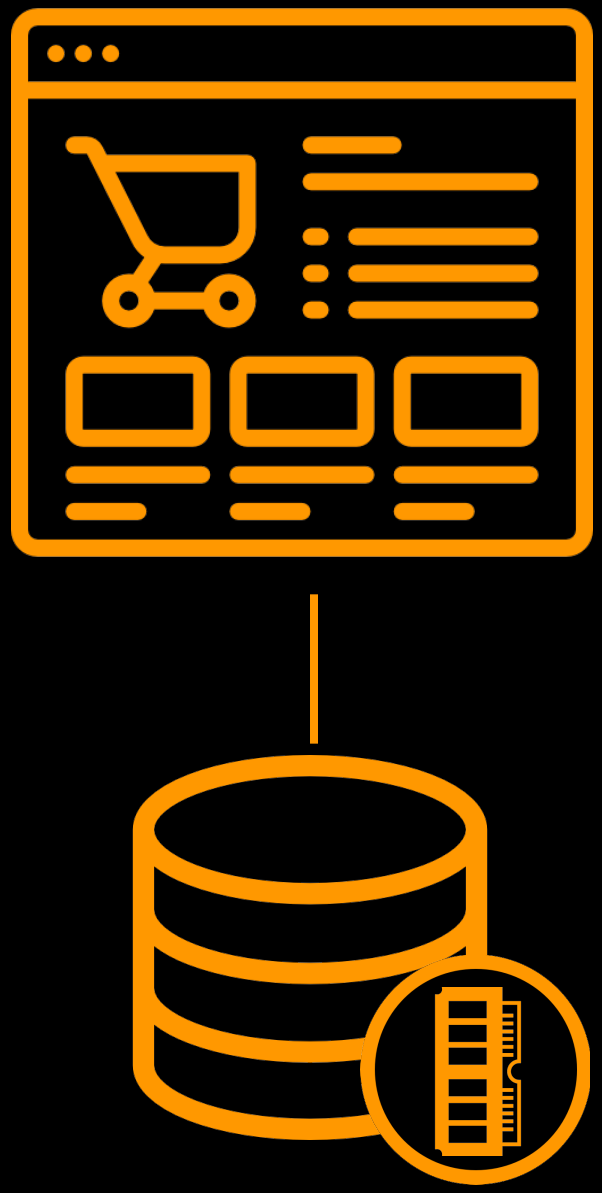

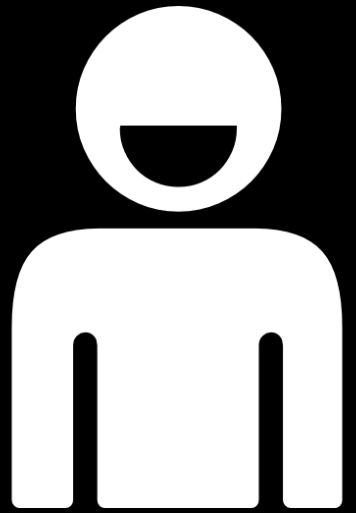

You are running out of memory

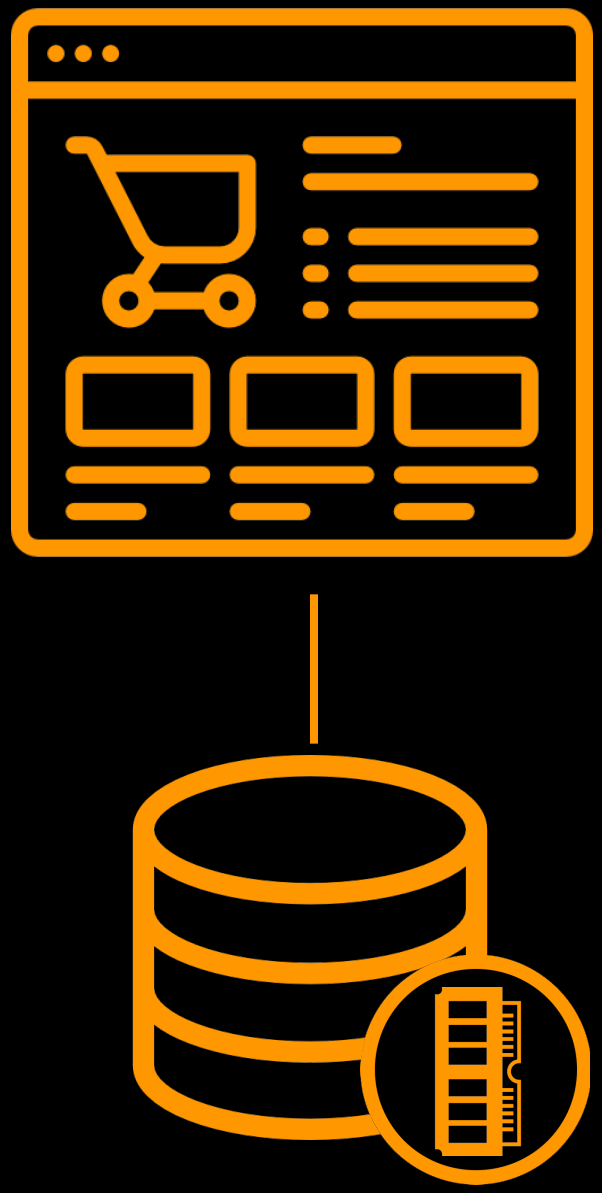

You are running out of memory

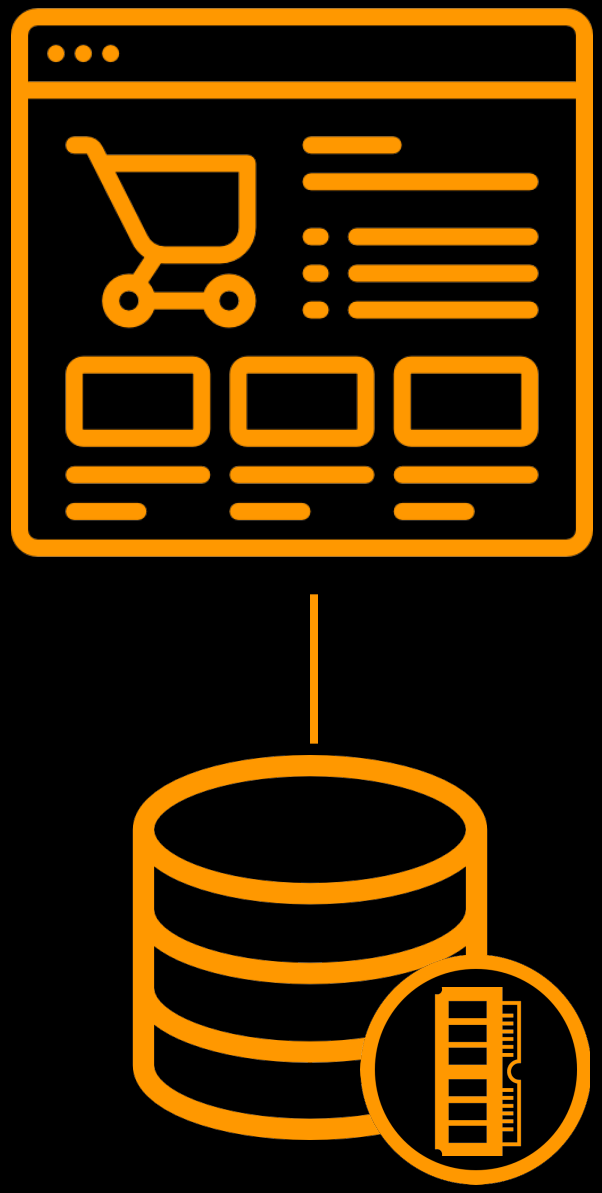

Buy more ?

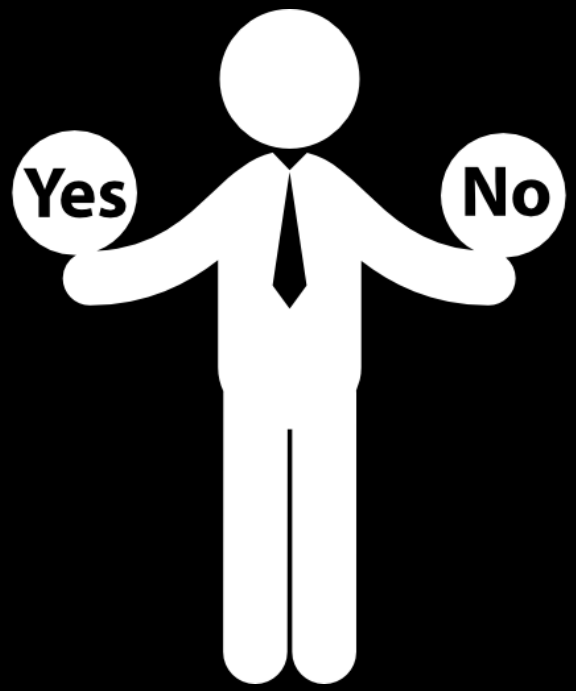

You are running out of memory

TPC-C on H-Store

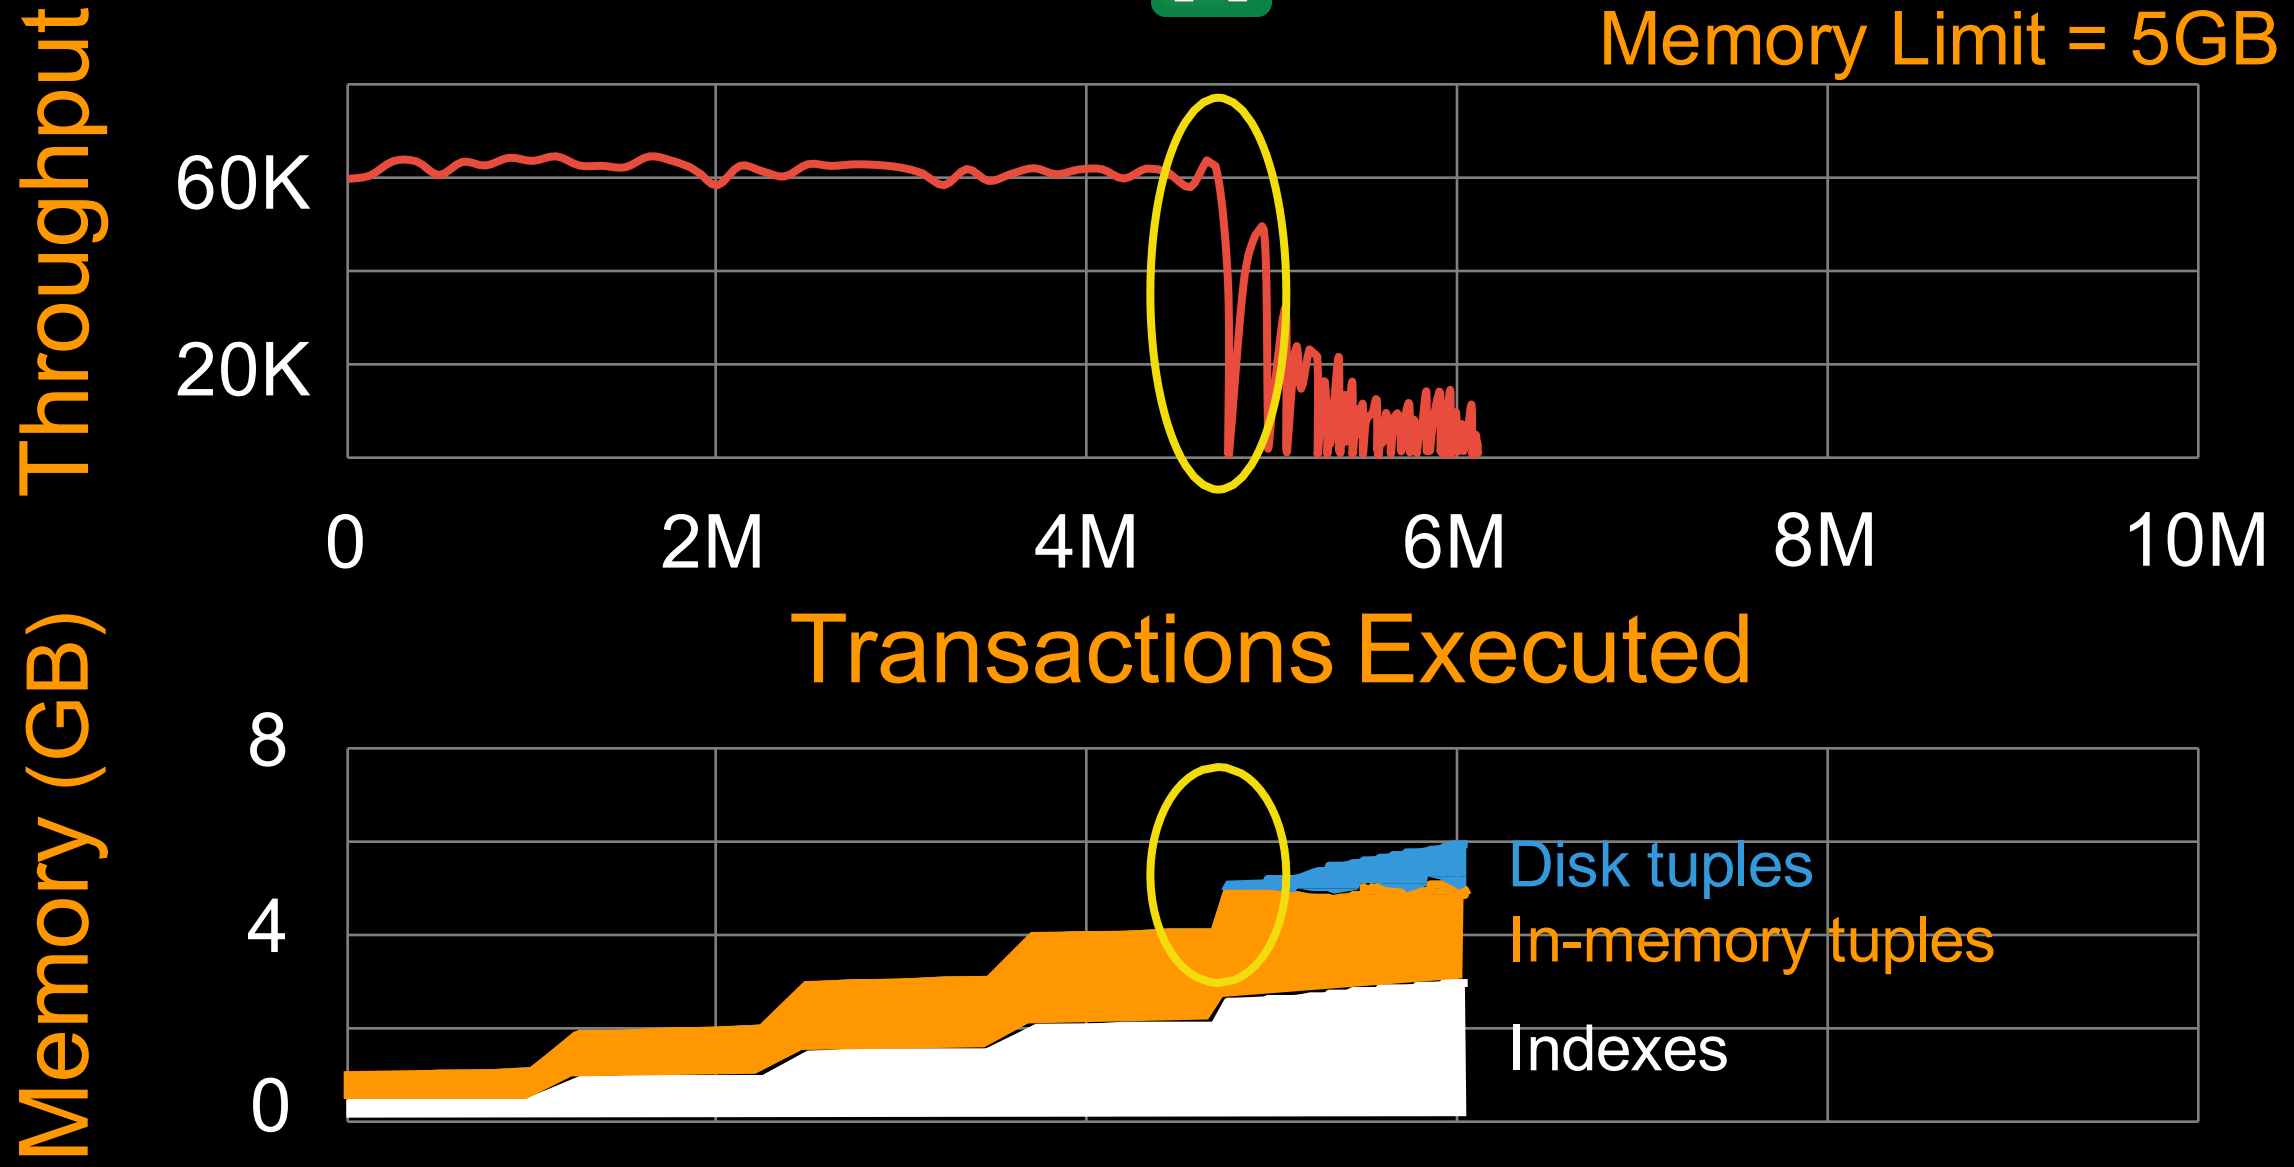


I GOT STUCK
SO I WENT TO SLEEP
zipmeme

The better way:

Use memory more efficiently

Indexes are LARGE

Benchmark	% space for index		Hybrid Index
TPC-C	58%	→	34%
Voter	55%	→	41%
Articles	34%	→	18%

Our Contributions

1. The hybrid index architecture
2. The Dual-Stage Transformation
3. Applied to 4 index structures

- B+tree
- Masstree
- Skip List
- Adaptive Radix Tree (ART)

Performance

Space

30 – 70%

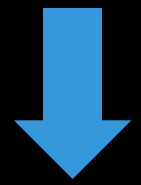

Did we solve this problem?

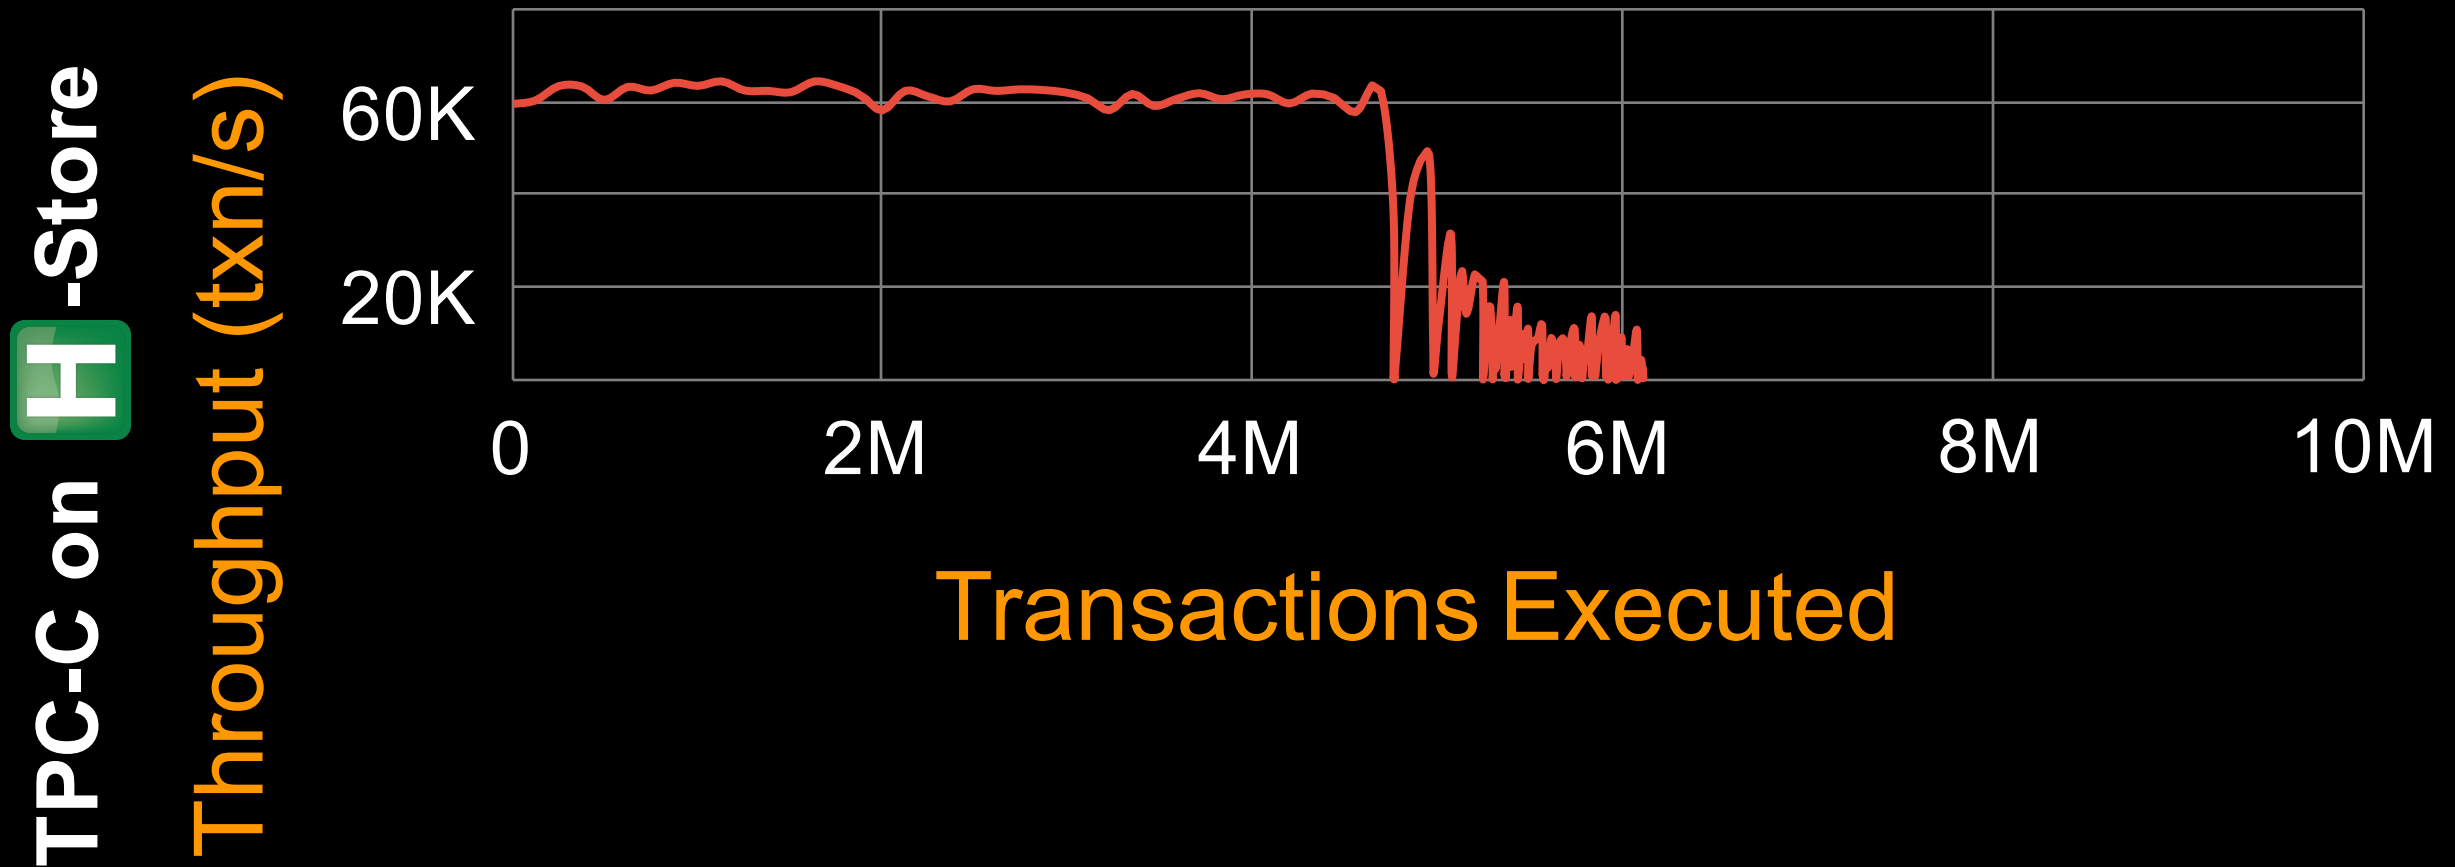

How do hybrid indexes achieve memory savings ?

🔑 Static

Hybrid Index: a dual-stage architecture

dynamic stage

static stage

Inserts are batched in the dynamic stage

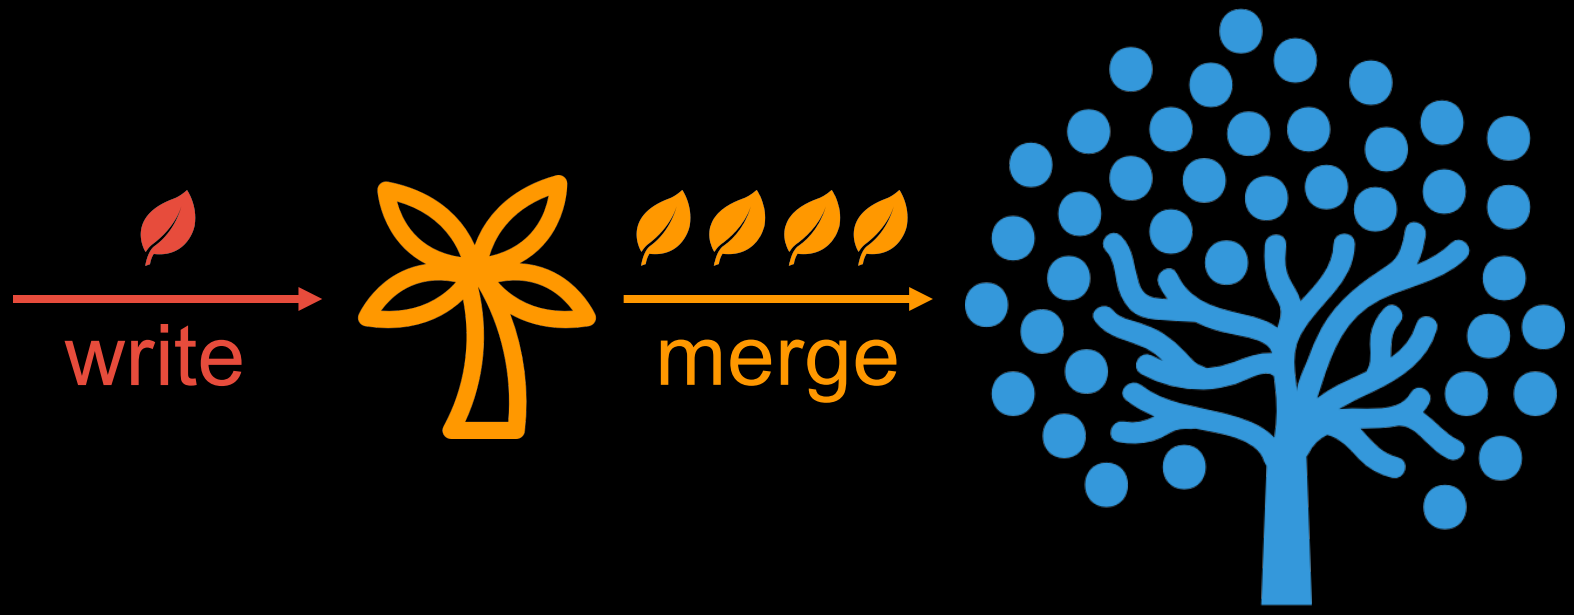

dynamic stage

static stage

Reads search the stages in order

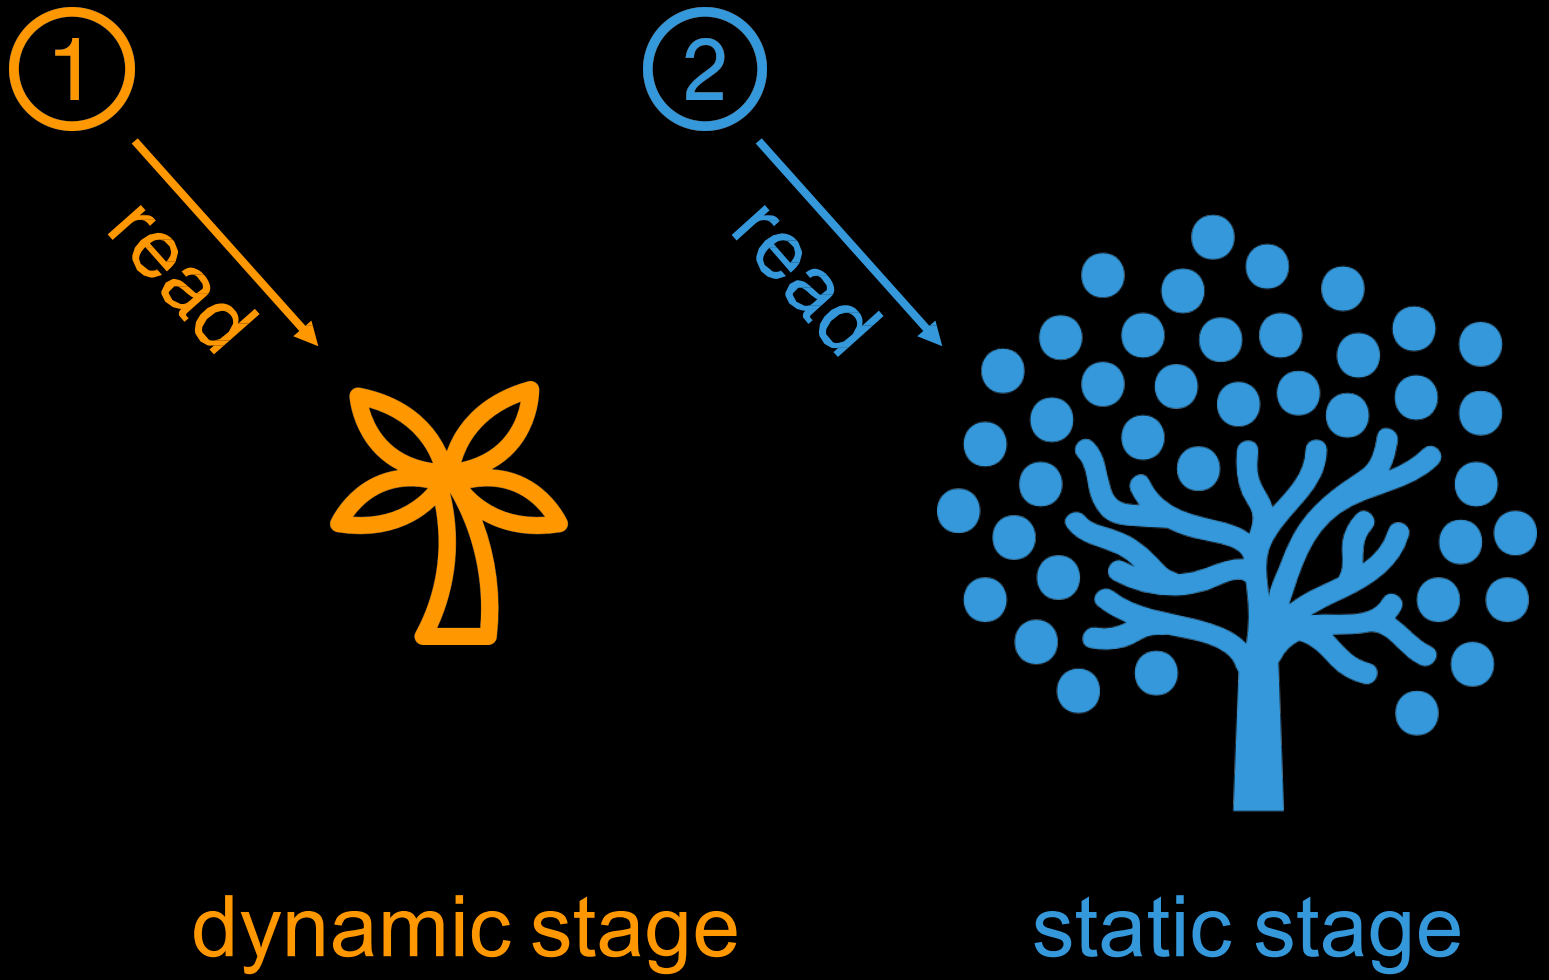

A Bloom filter improves read performance

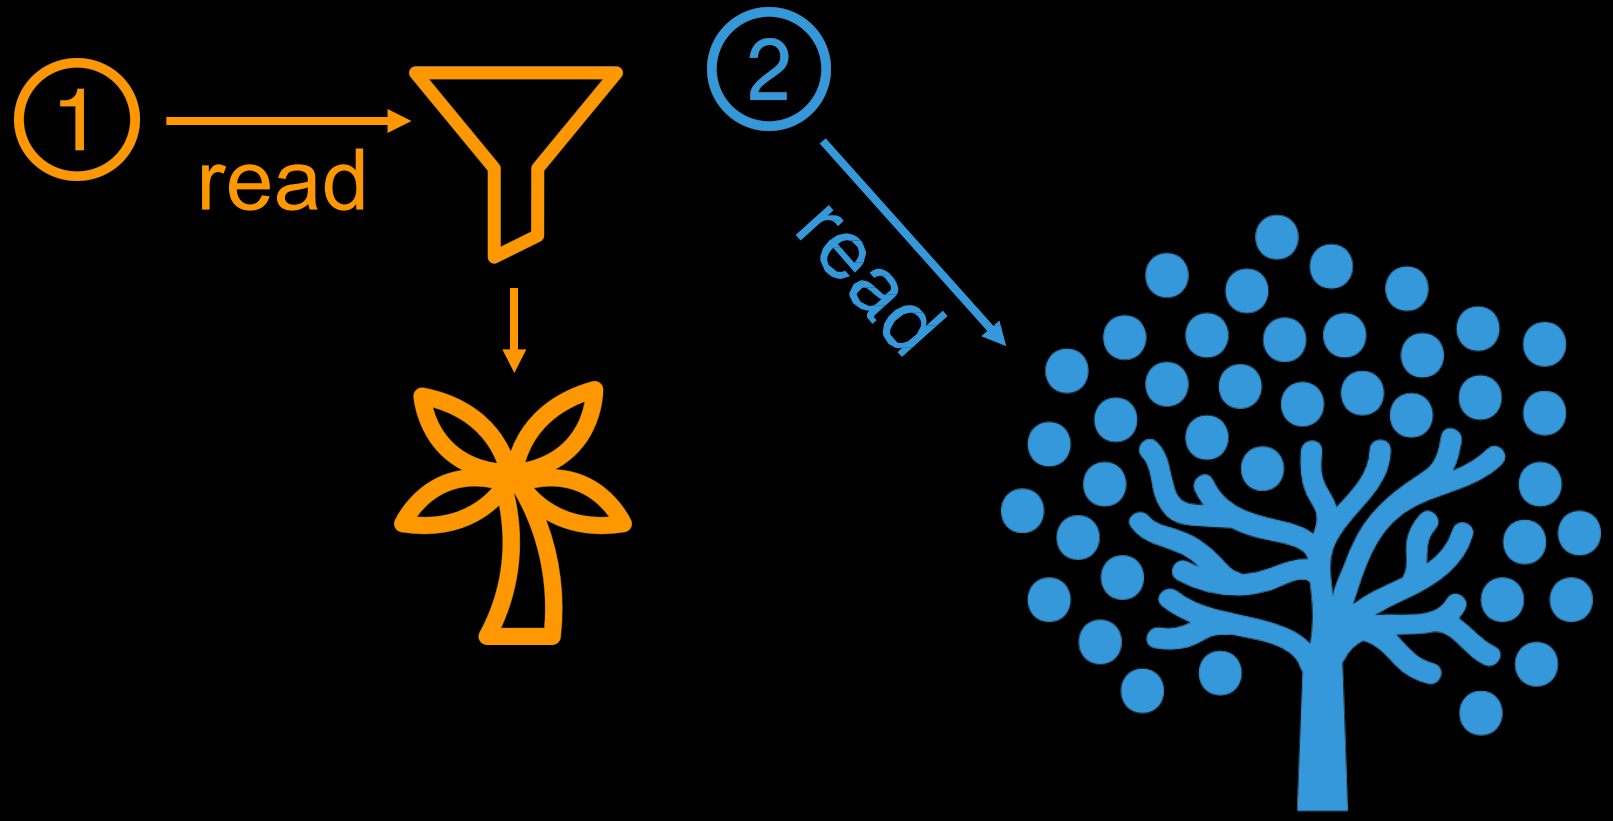

dynamic stage

static stage

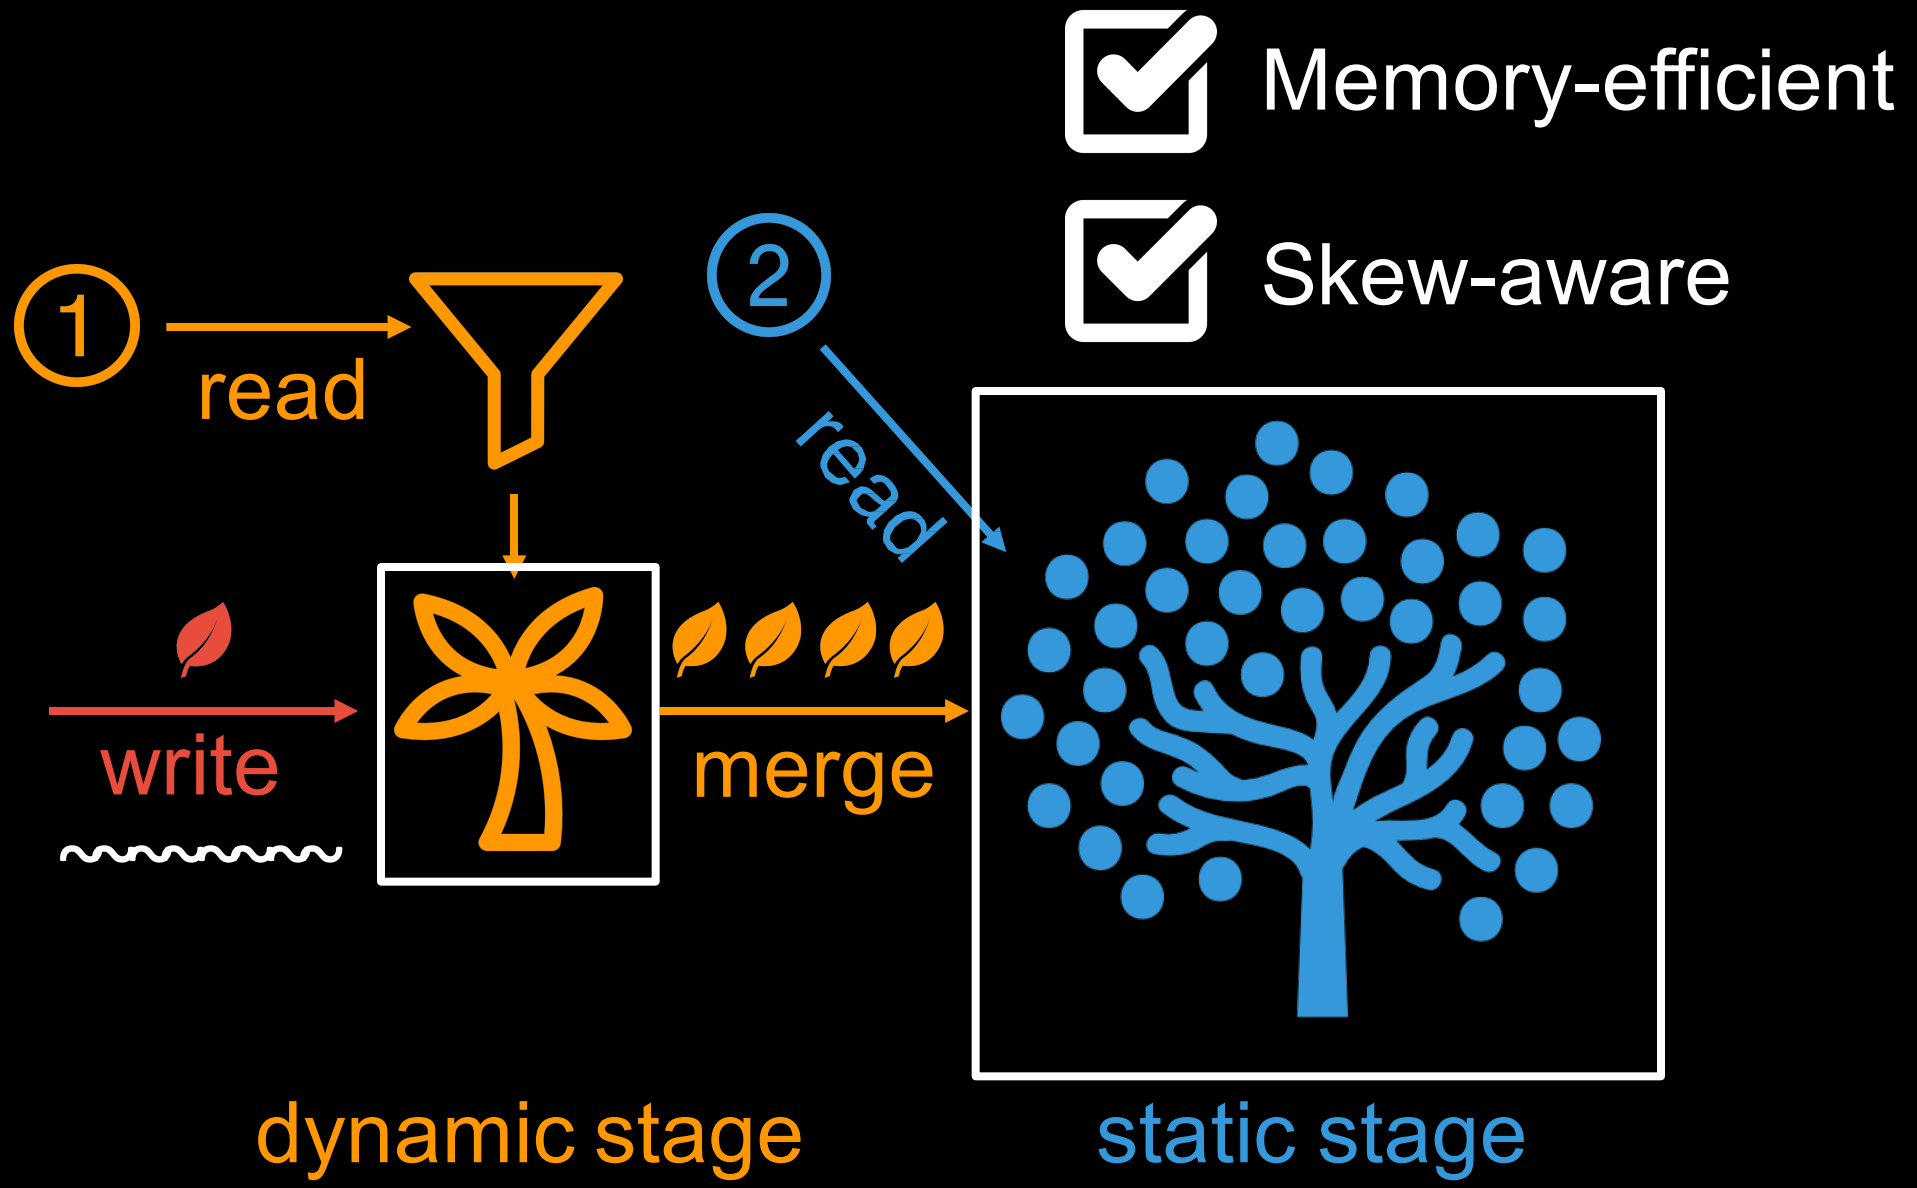
Memory-efficient
Skew-aware
1
read
2
read
write
merge
dynamic stage
static stage

The Dual-Stage Transformation

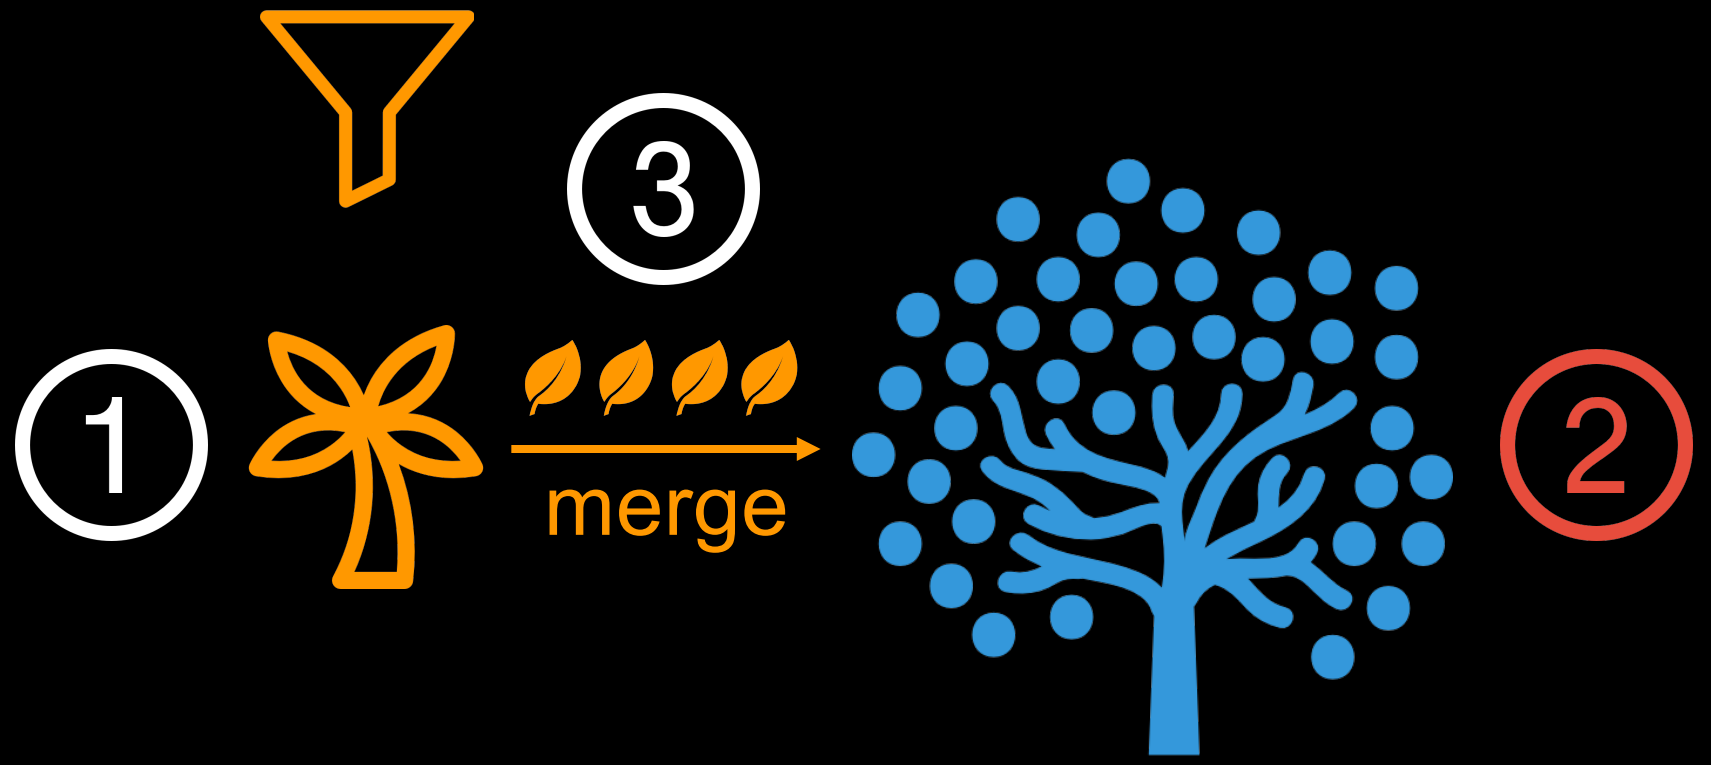

The Dynamic-to-Static Rules

Compaction

Reduction

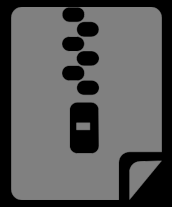

Compression

Compaction: minimize # of memory blocks

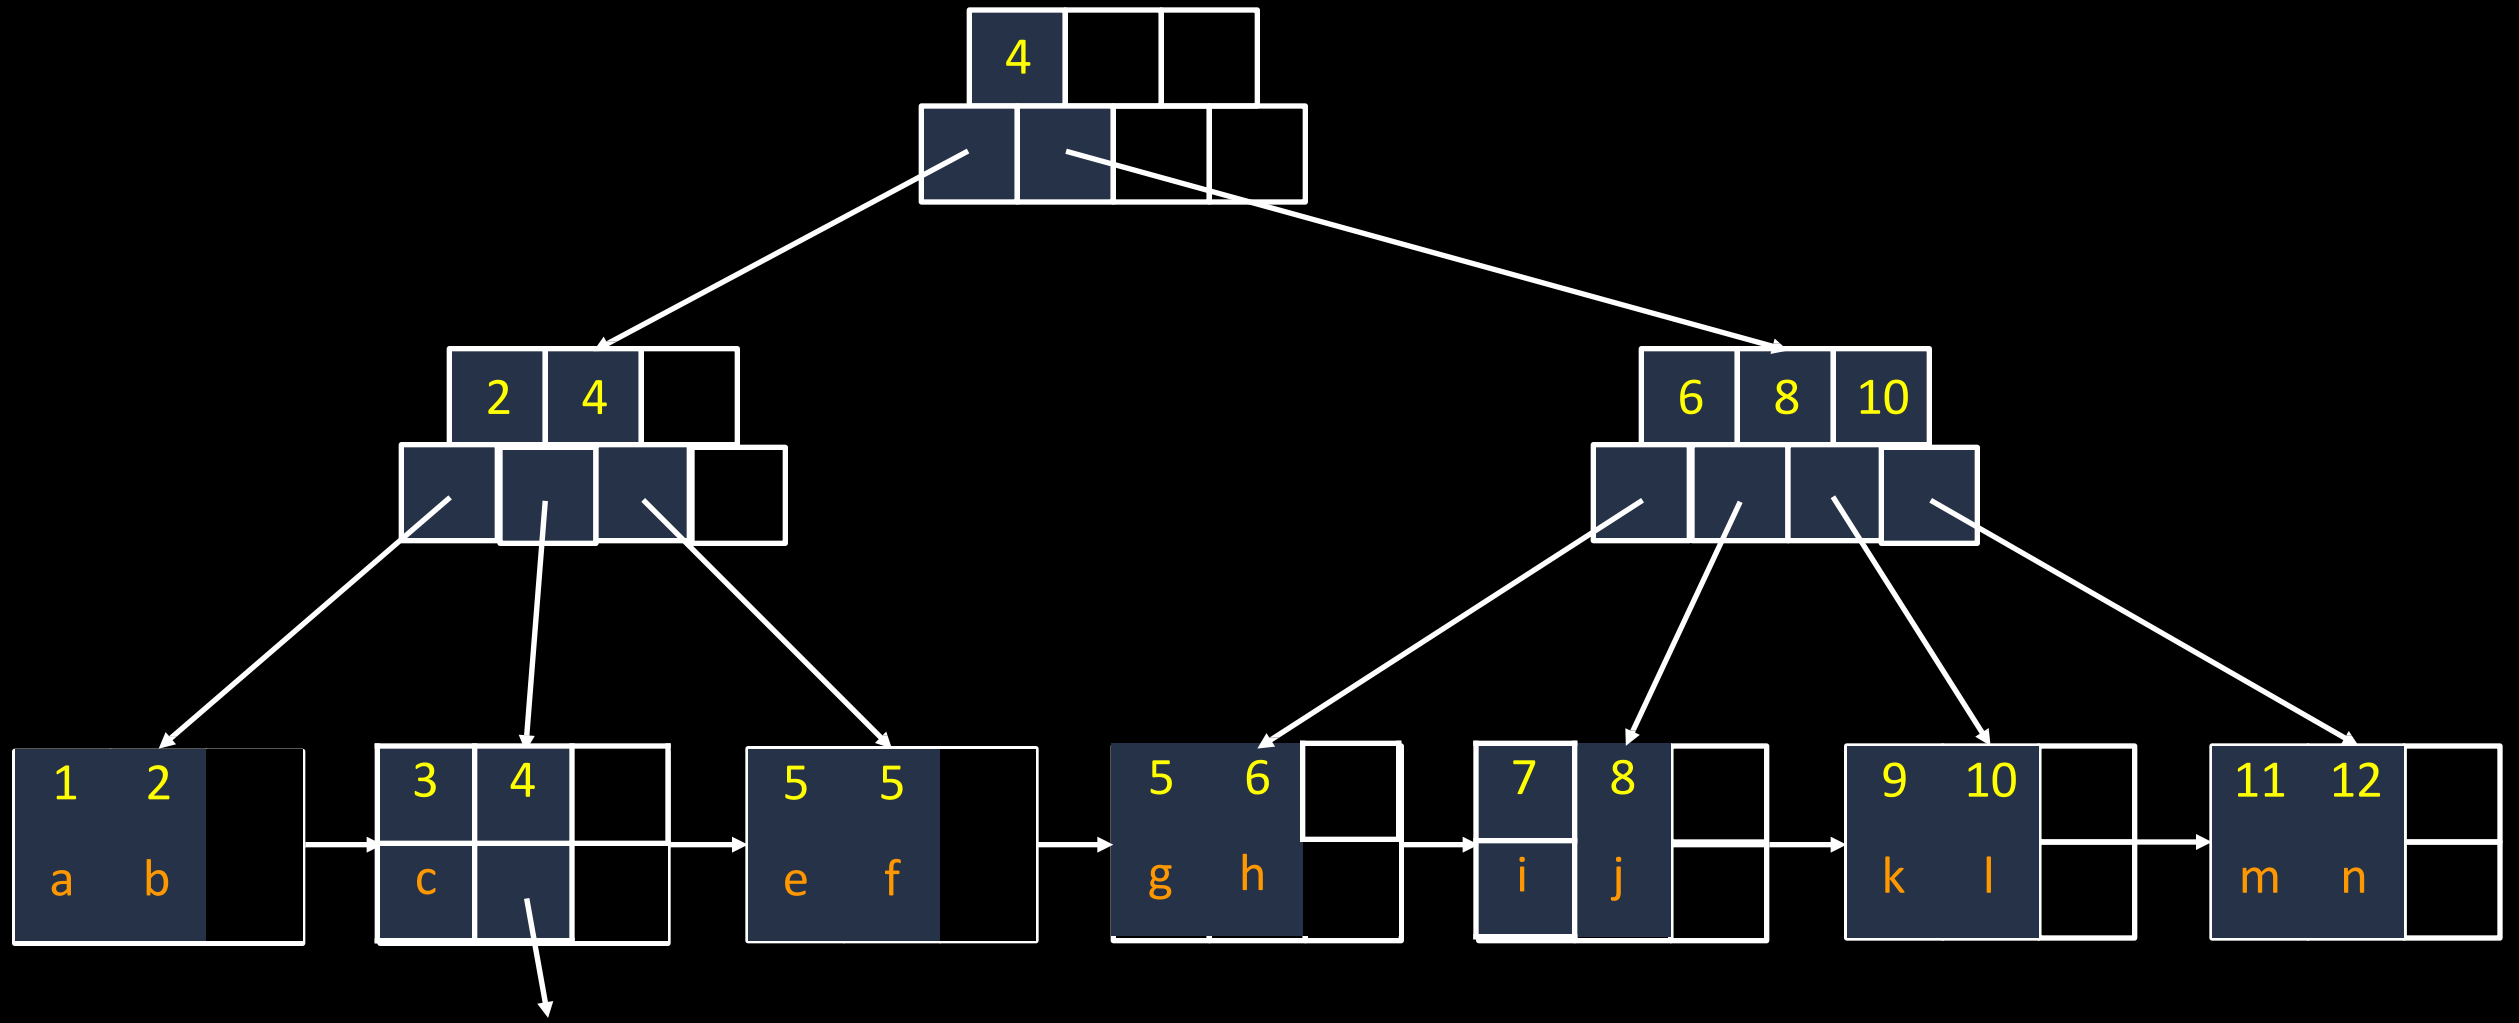

Compaction: minimize # of memory blocks

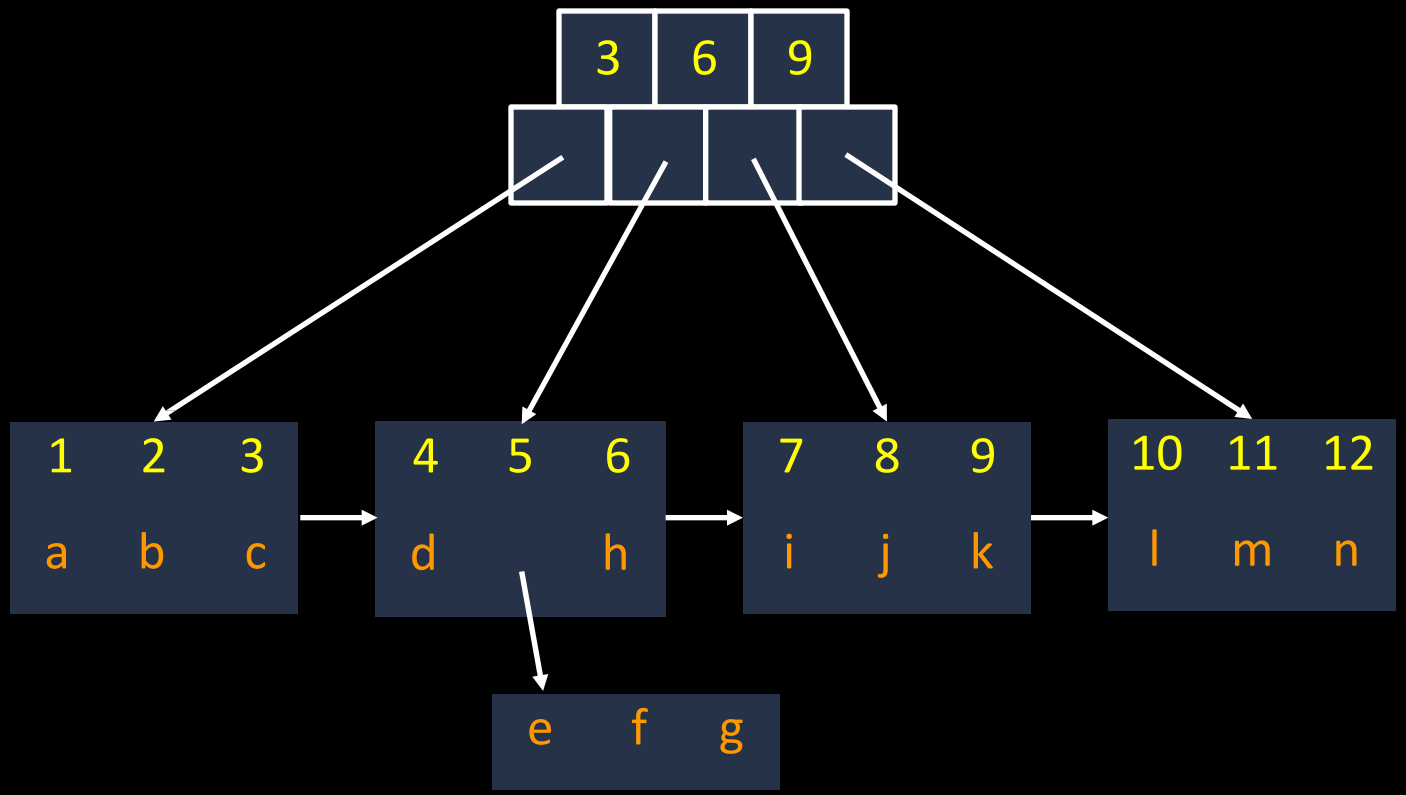

Reduction: minimize structural overhead

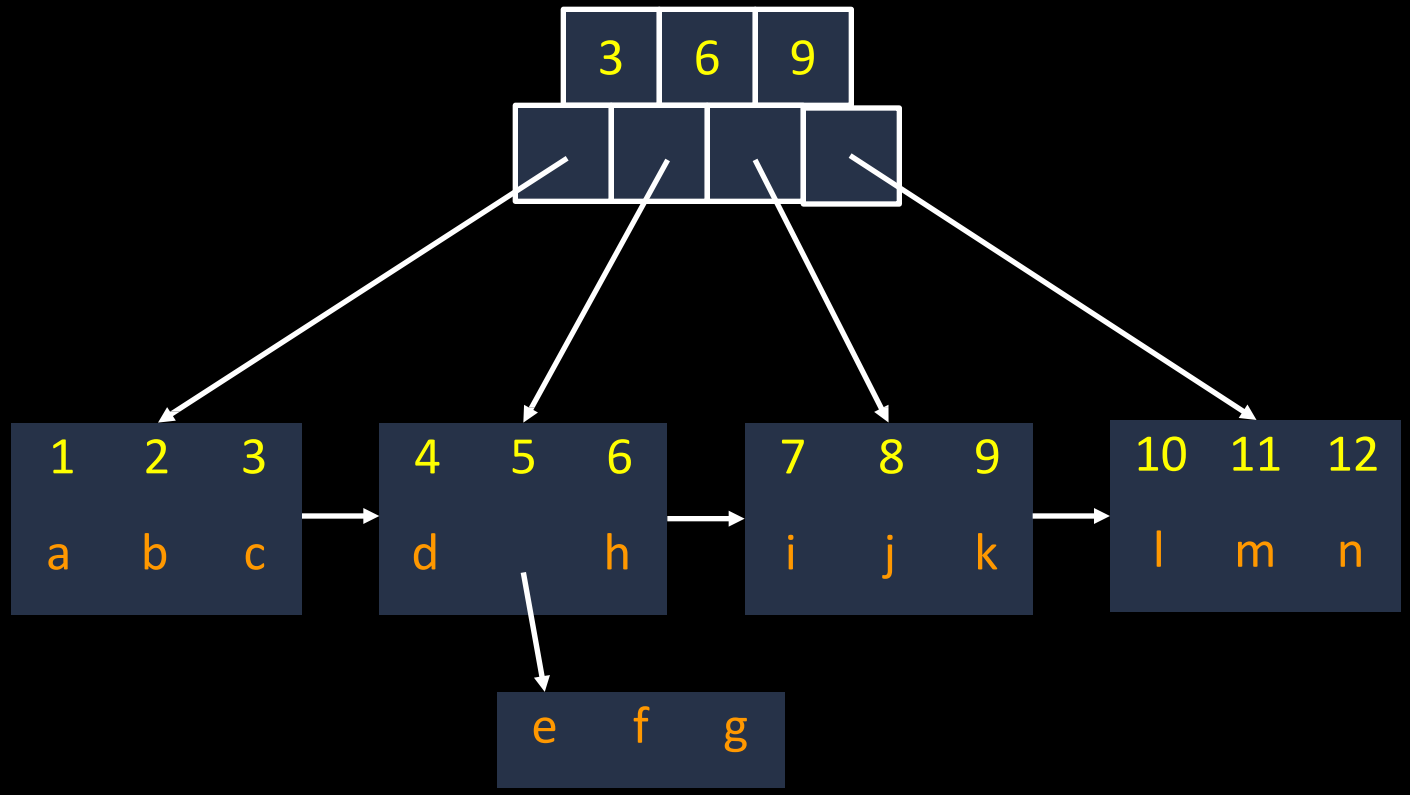

Reduction: minimize structural overhead

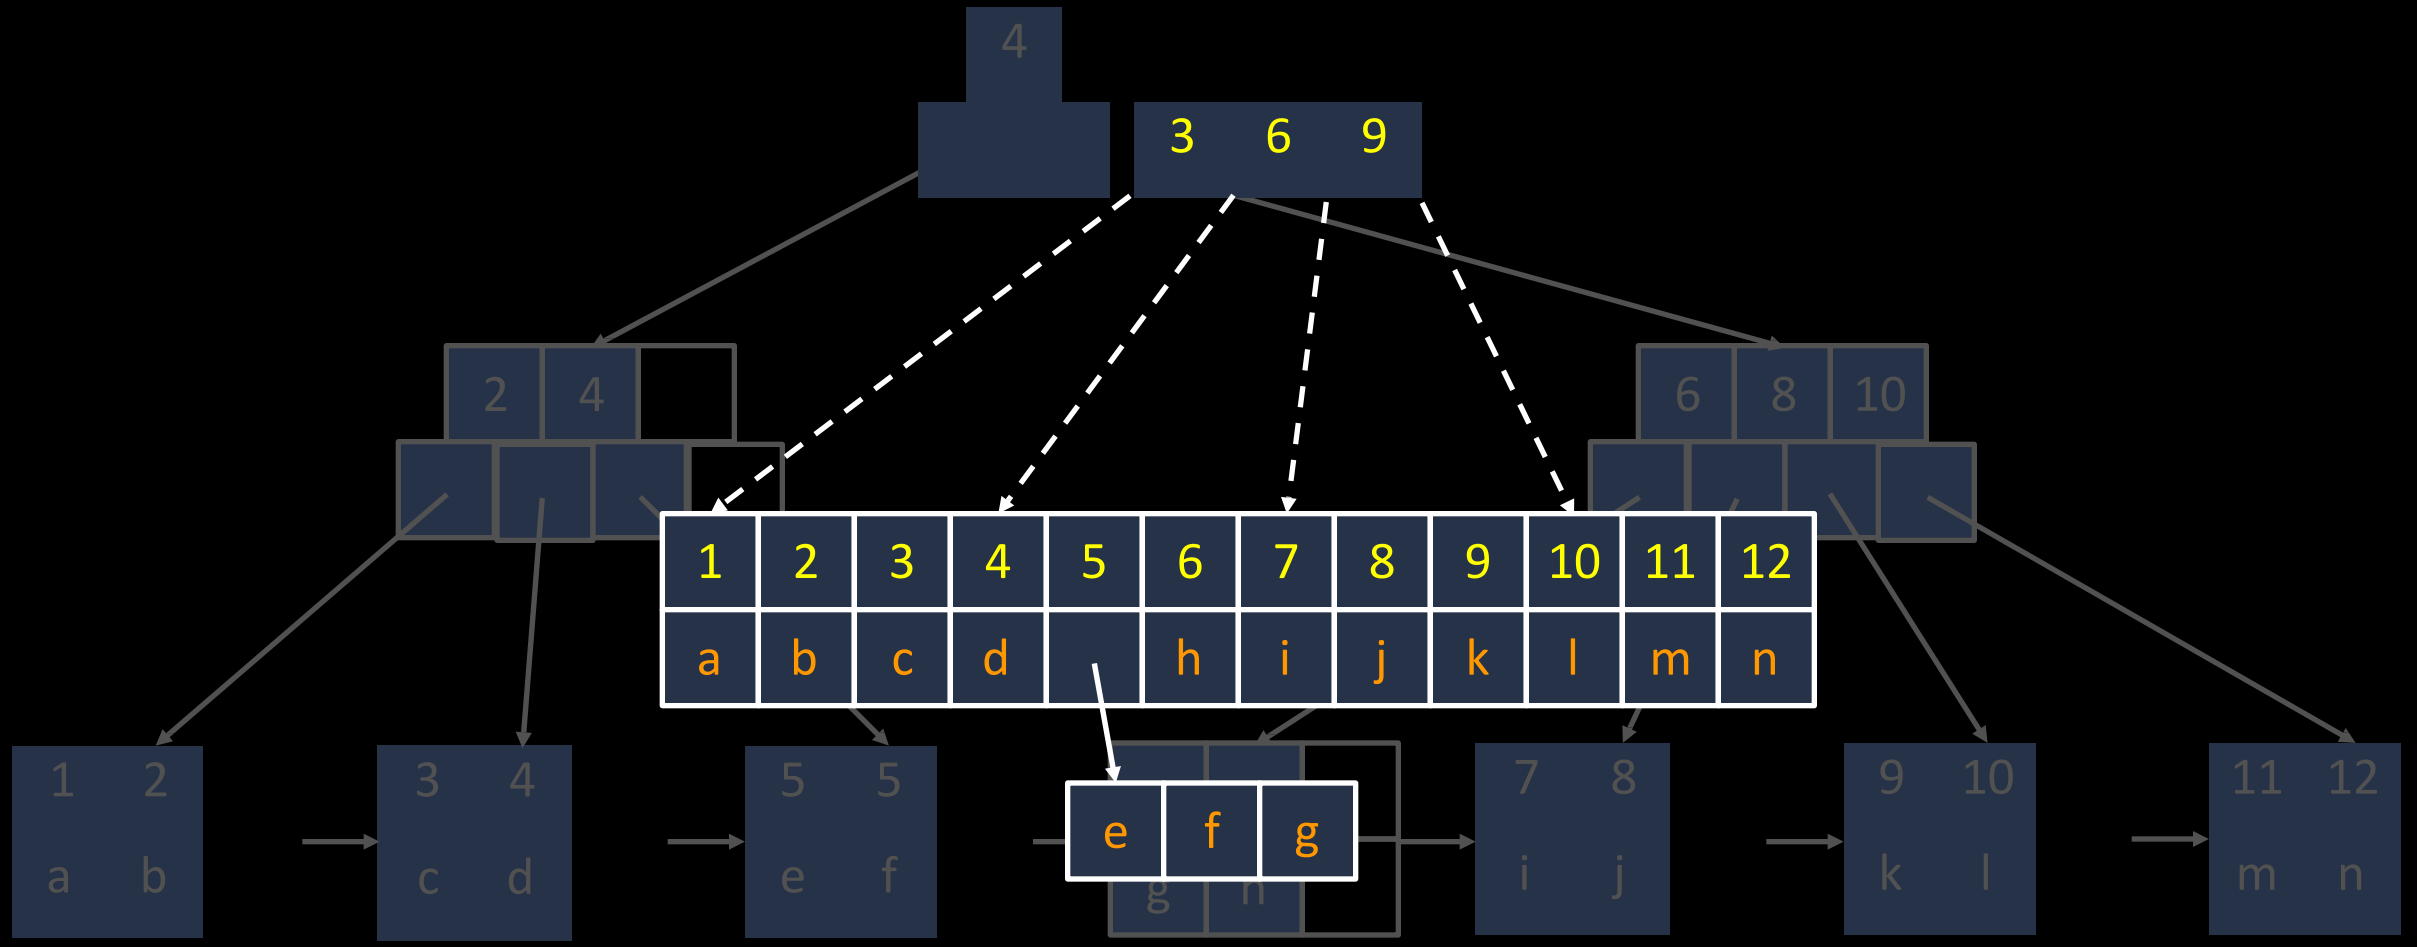

The Dual-Stage Transformation

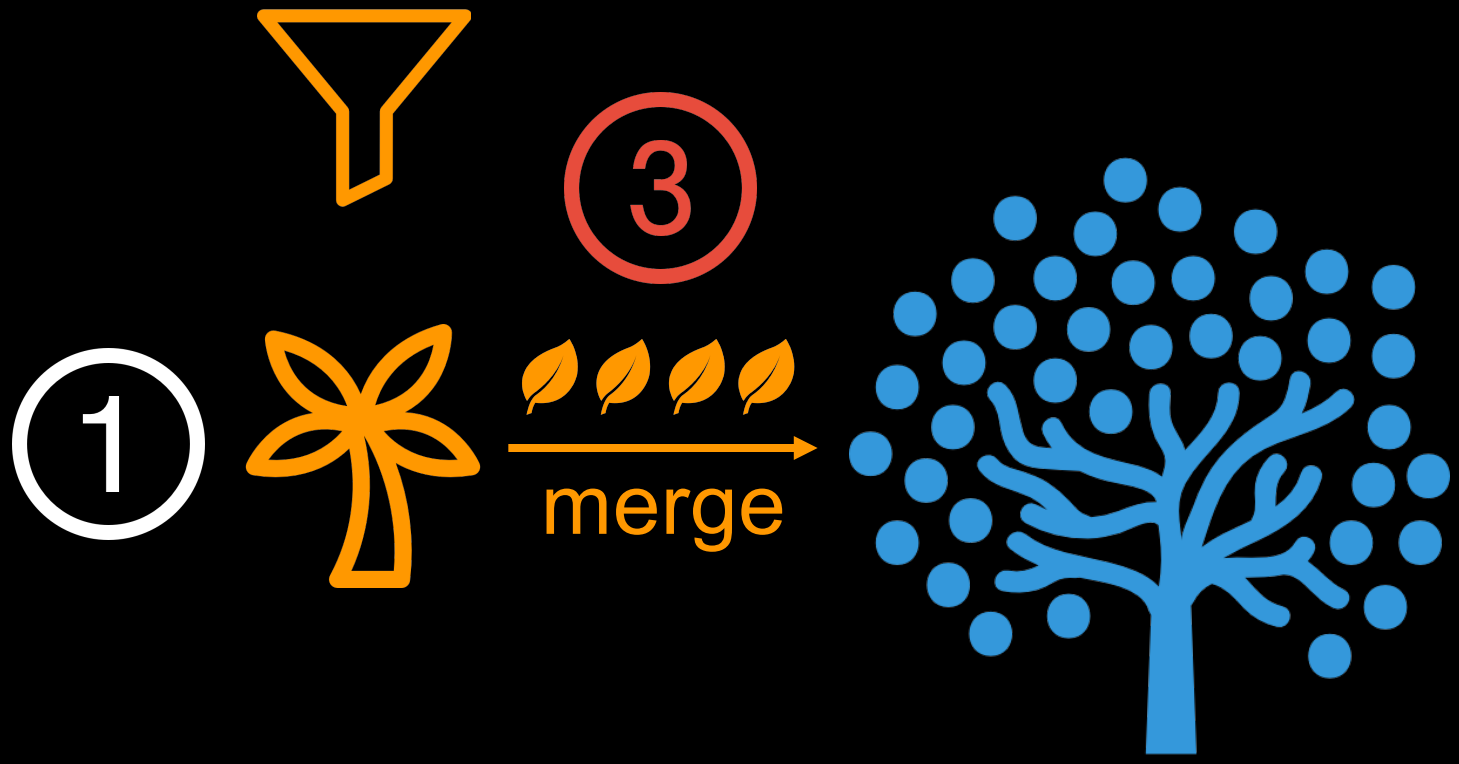

dynamic stage

static stage

The Dual-Stage Transformation

Merge Questions:

1. Partial?
2. When?
3. Blocking?

dynamic stage

static stage

Did we solve this problem?

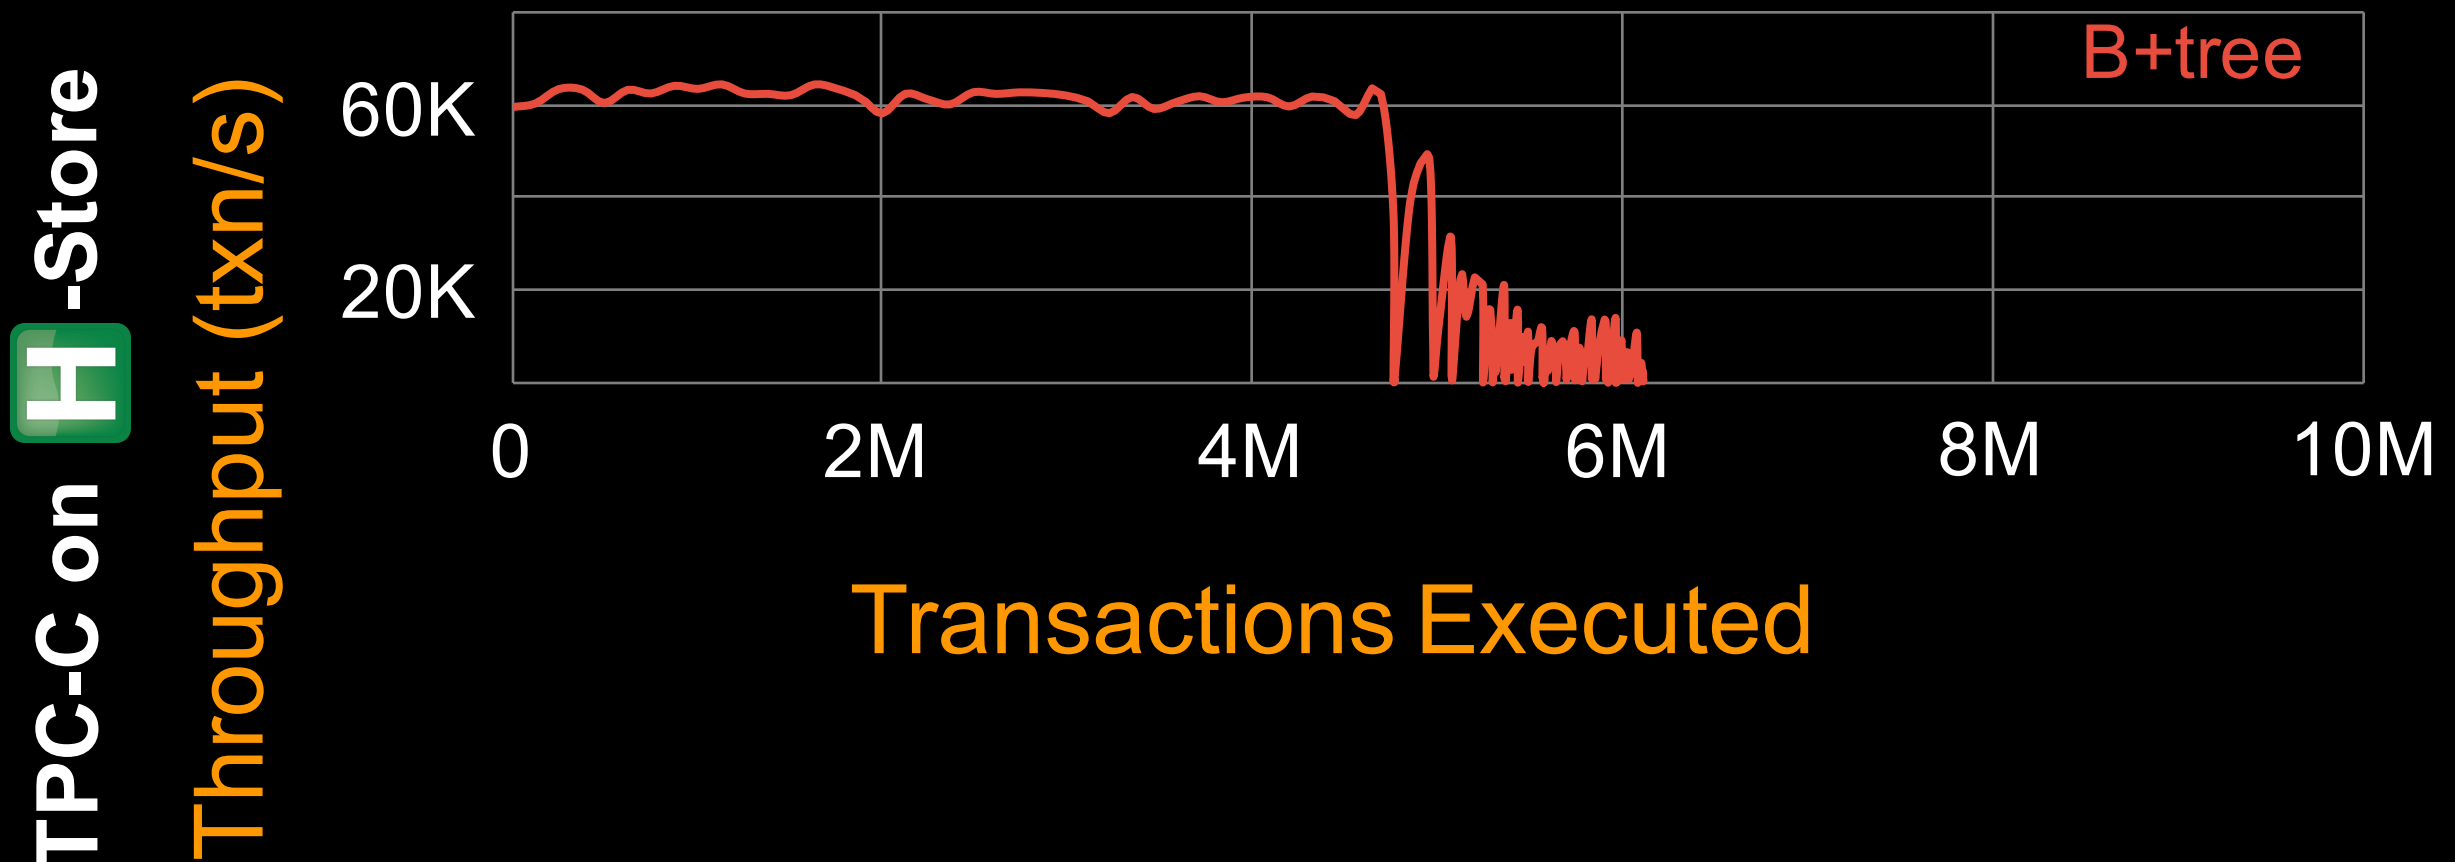

Yes, we improved the DBMS's capacity!

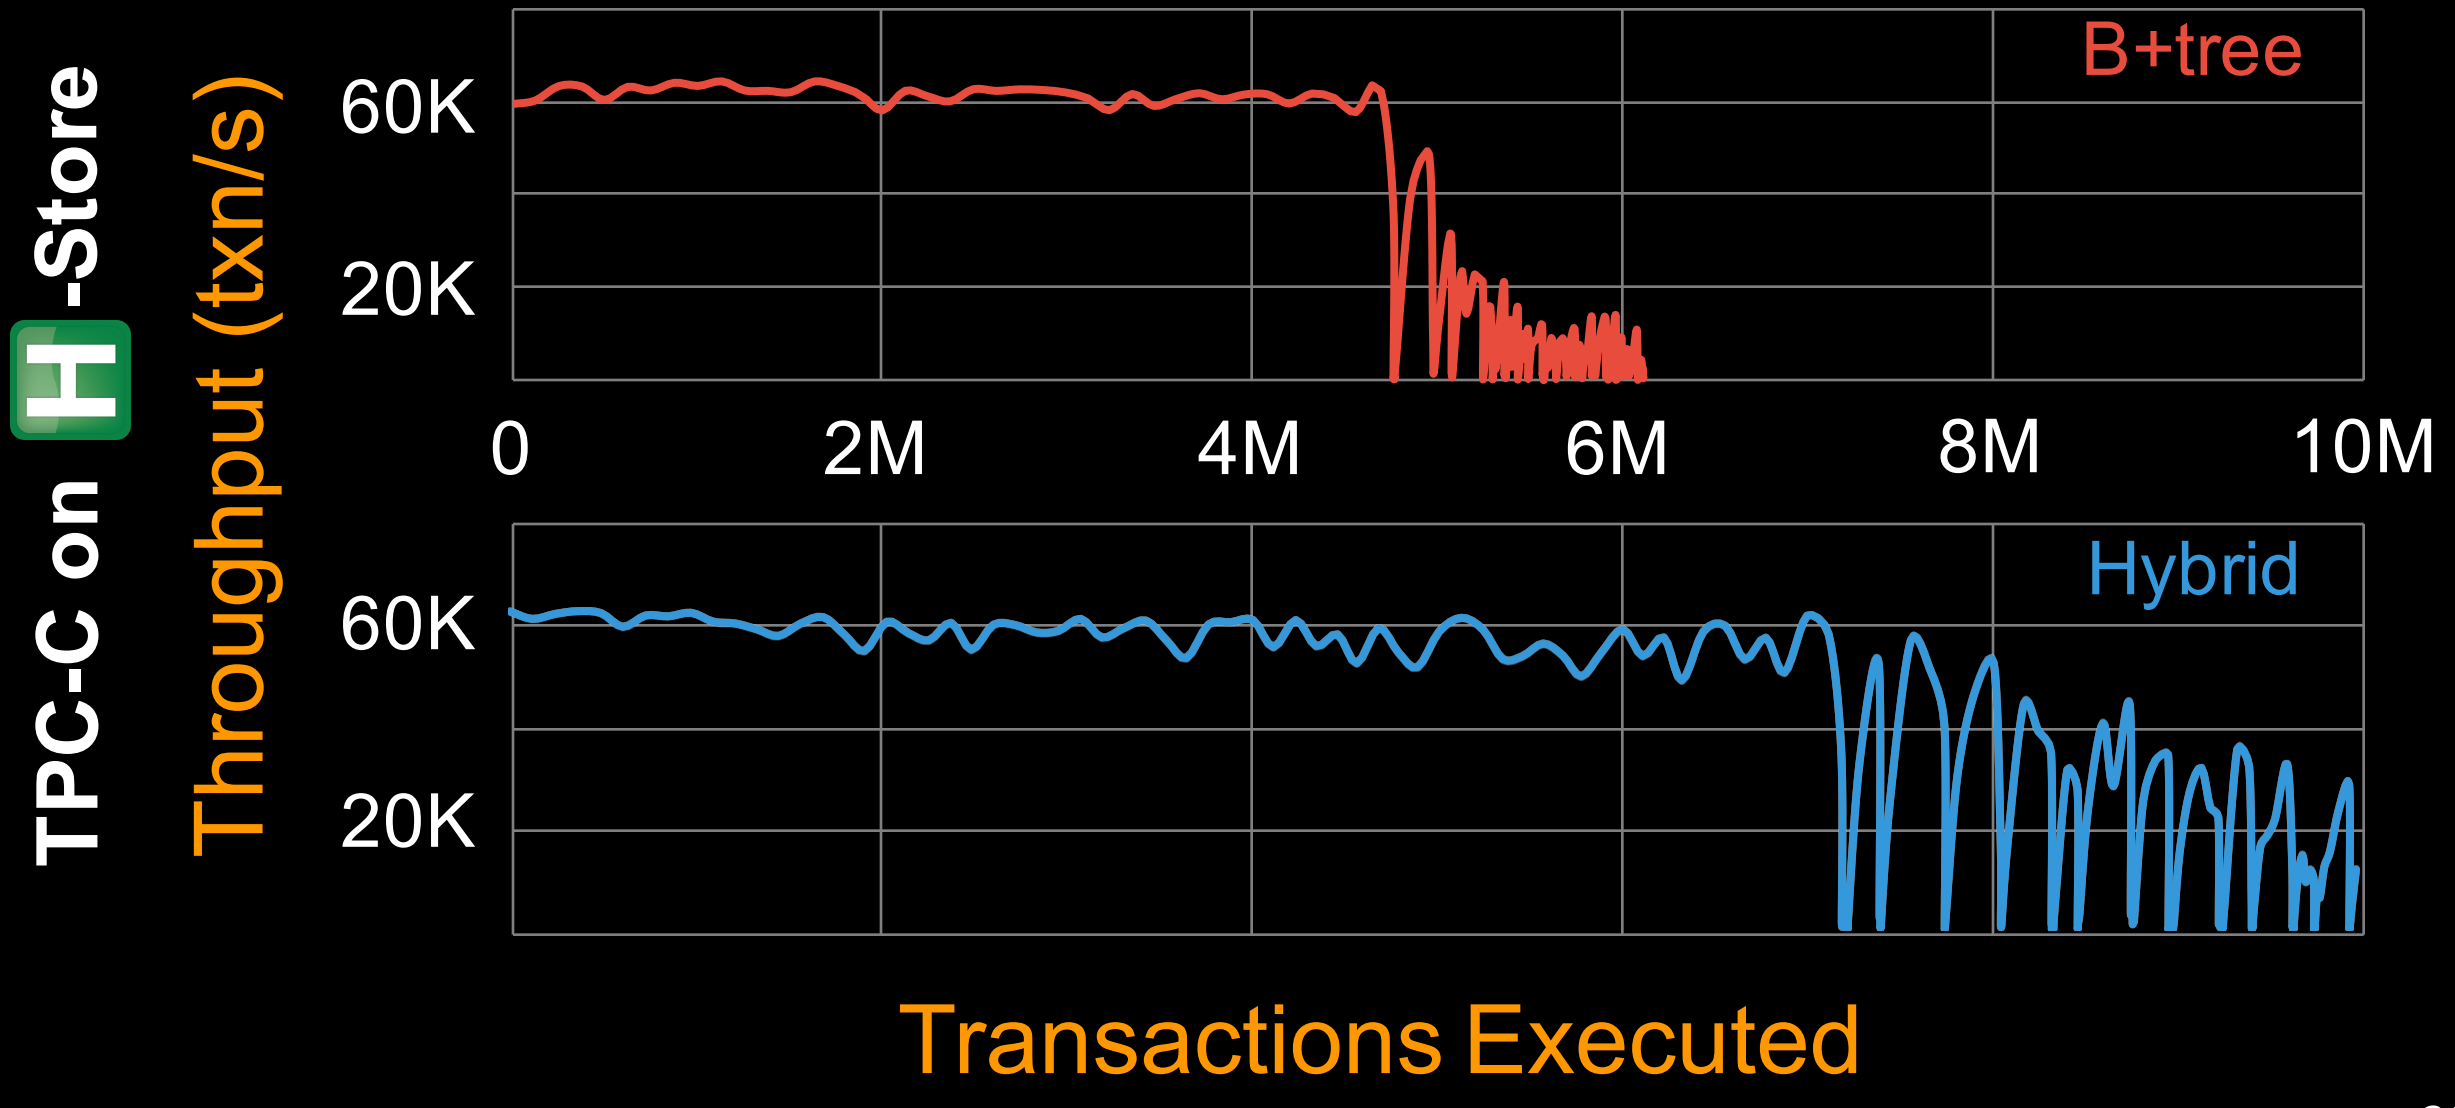

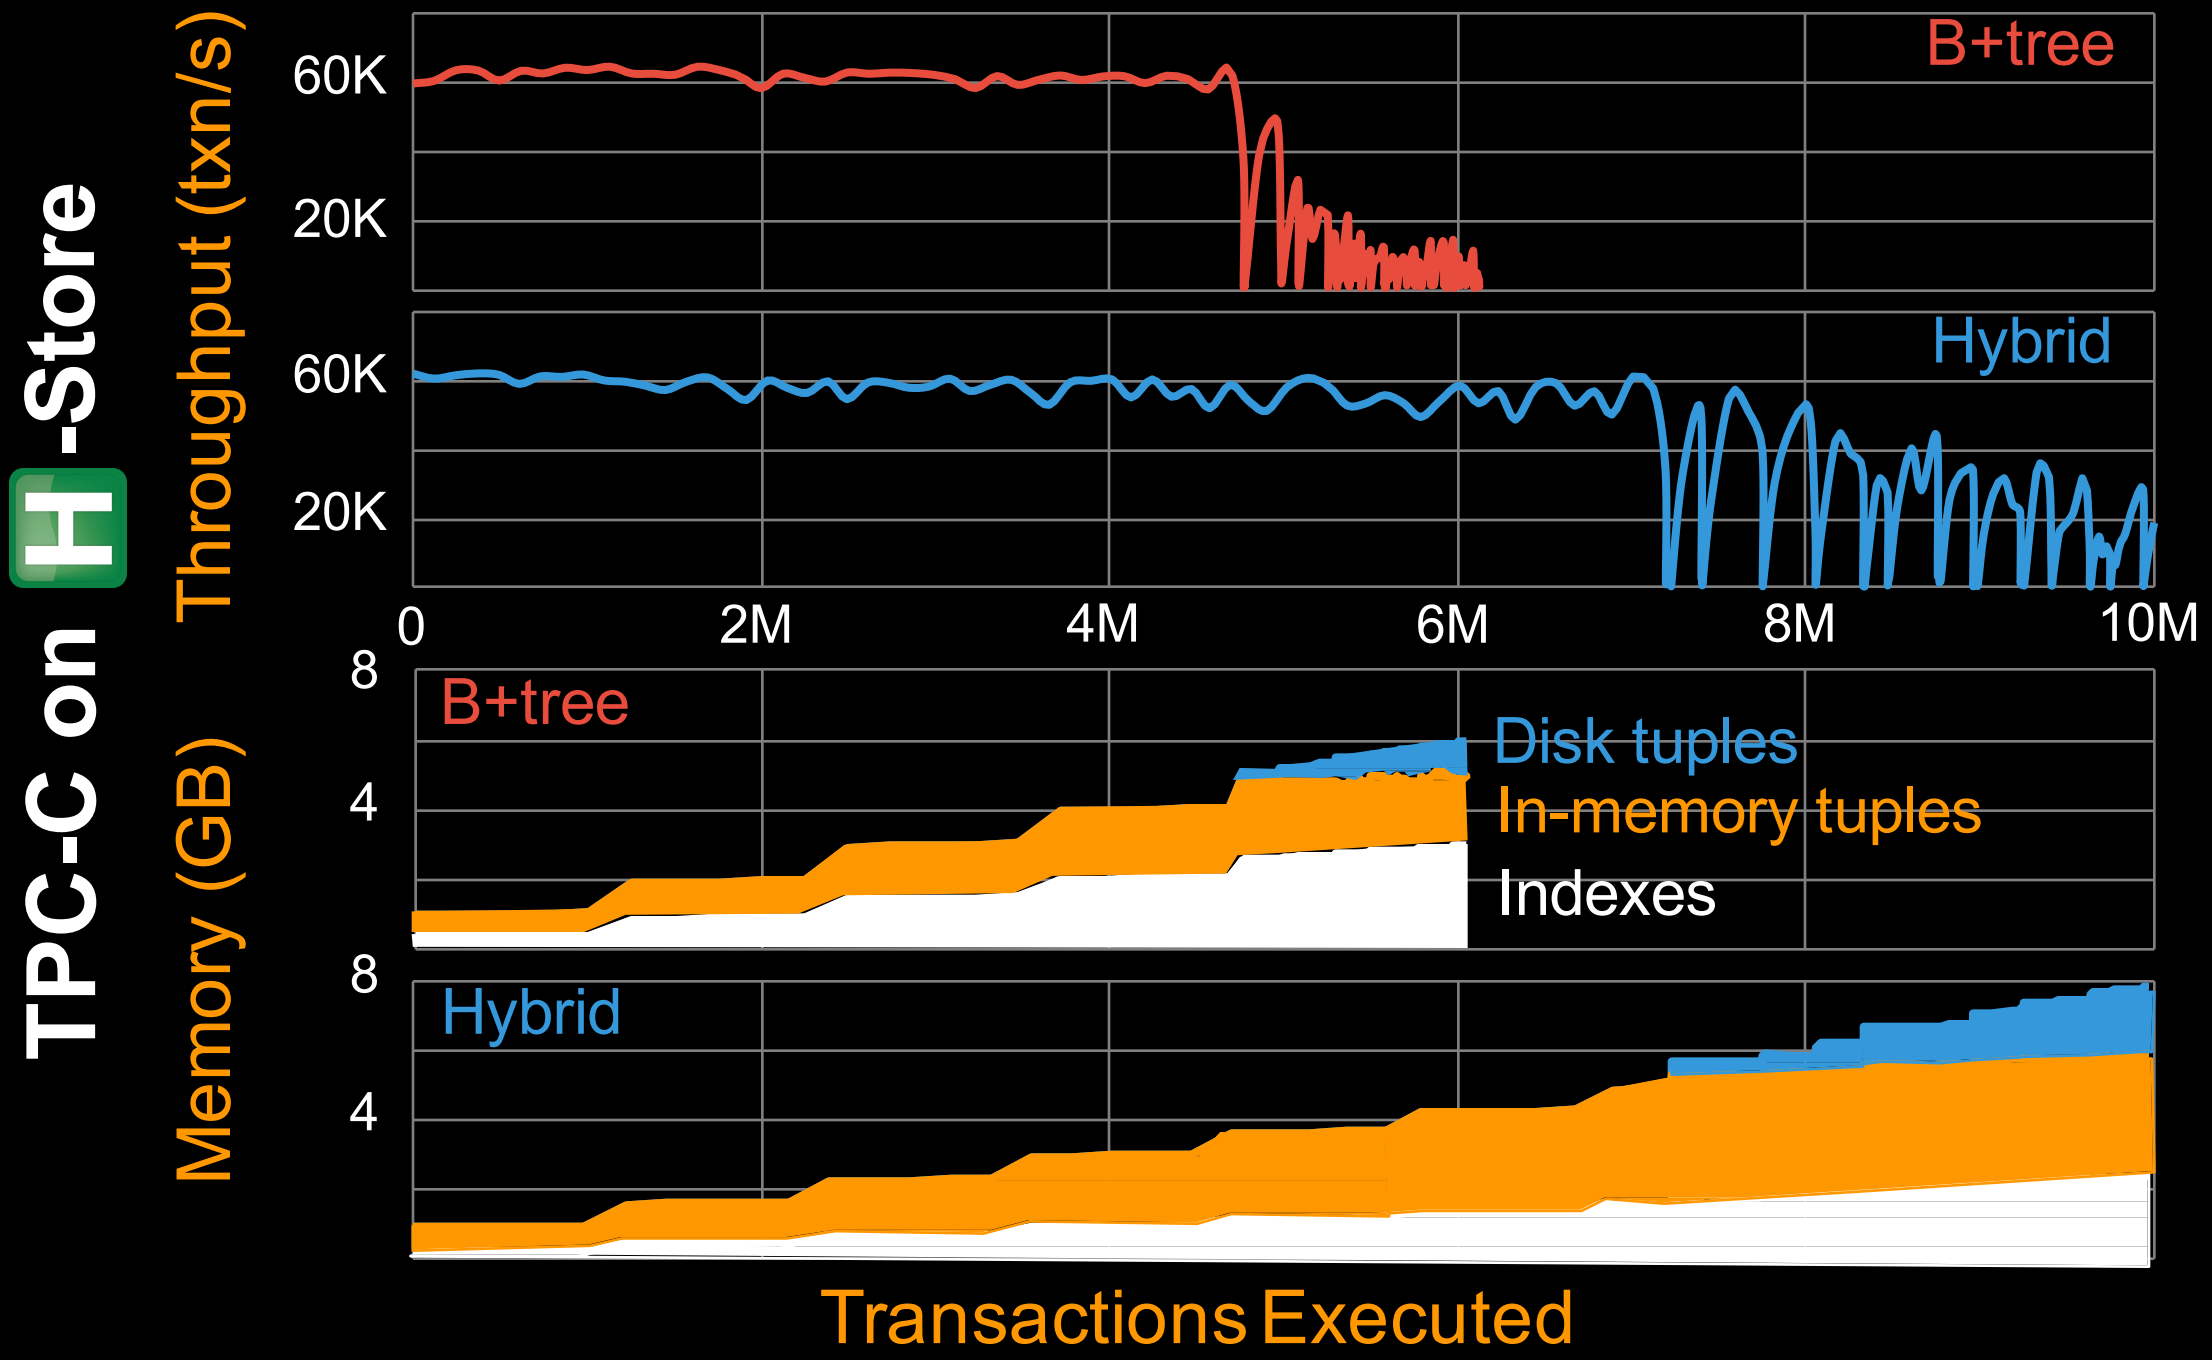
TPC-C on H-Store
Throughput (txn/s)
60K
20K
B+tree
60K
20K
Hybrid
0
2M
4M
6M
8M
10M
Memory (GB)
8
4
B+tree
Disk tuples
In-memory tuples
Indexes
8
4
Hybrid
Transactions Executed

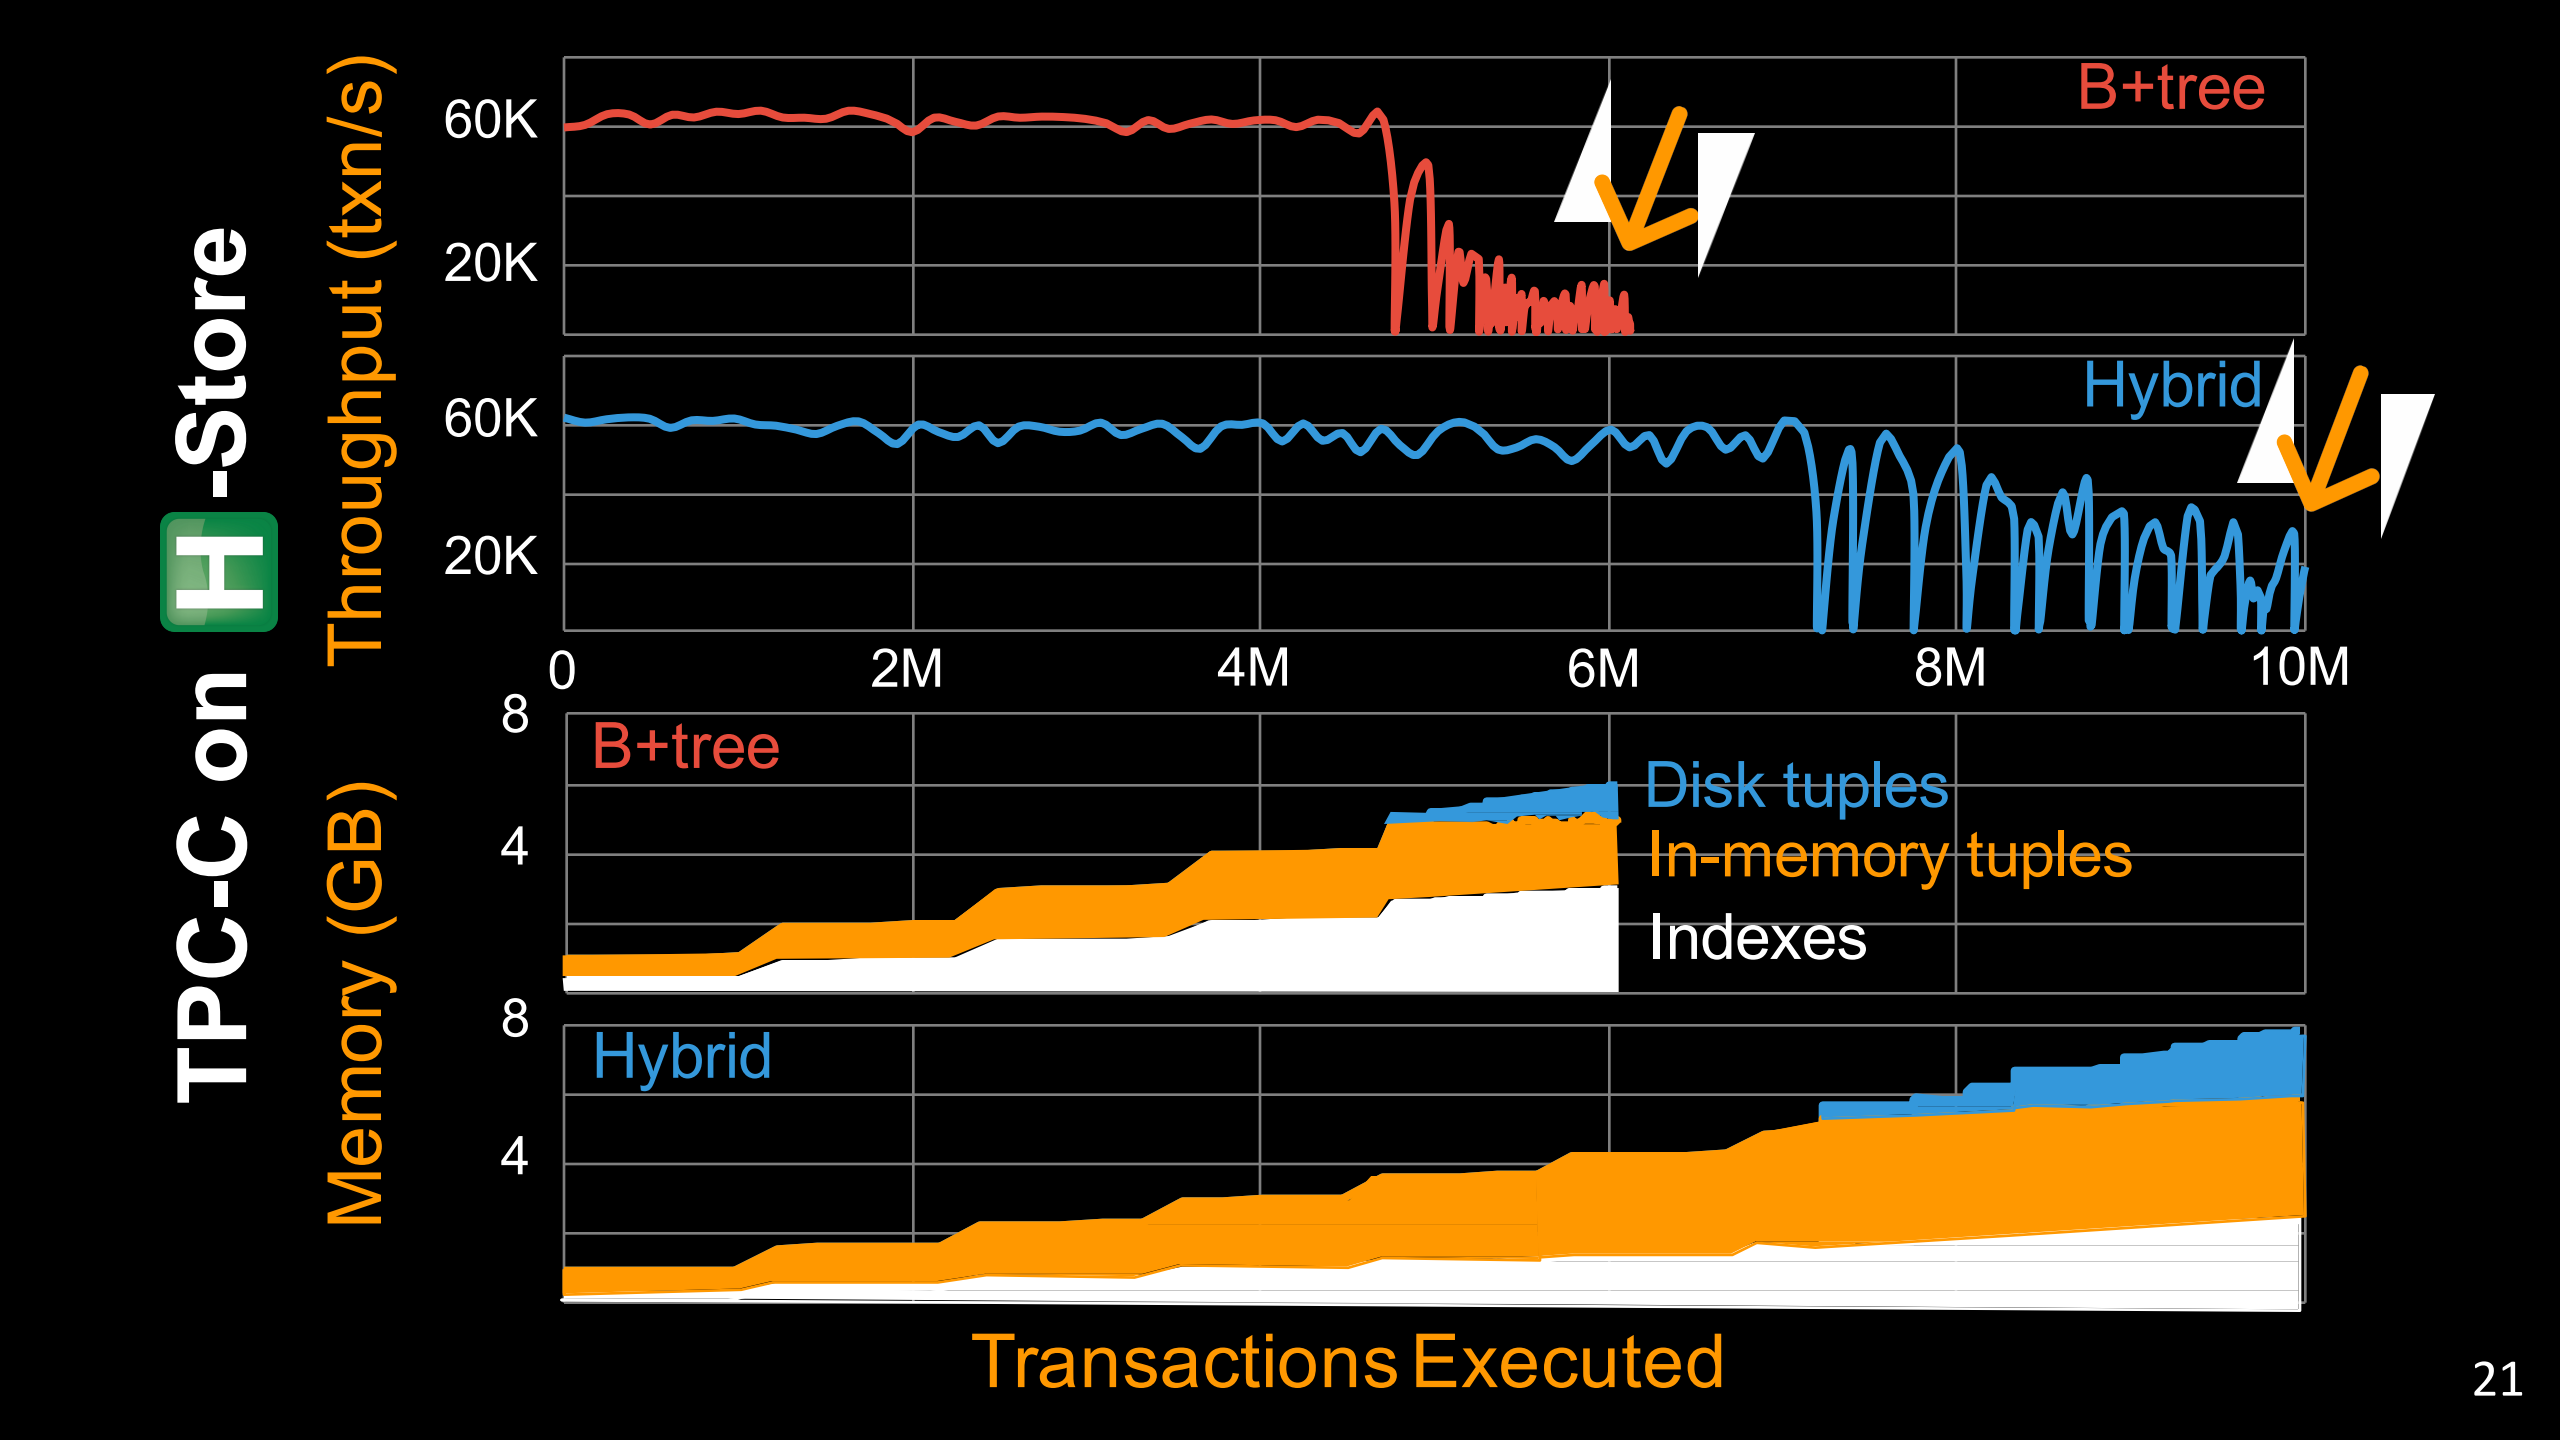
TPC-C on H-Store
Throughput (txn/s)
60K
20K
B+tree
60K
20K
Hybrid
0
2M
4M
6M
8M
10M
Memory (GB)
8
4
B+tree
Disk tuples
In-memory tuples
Indexes
8
4
Hybrid
Transactions Executed

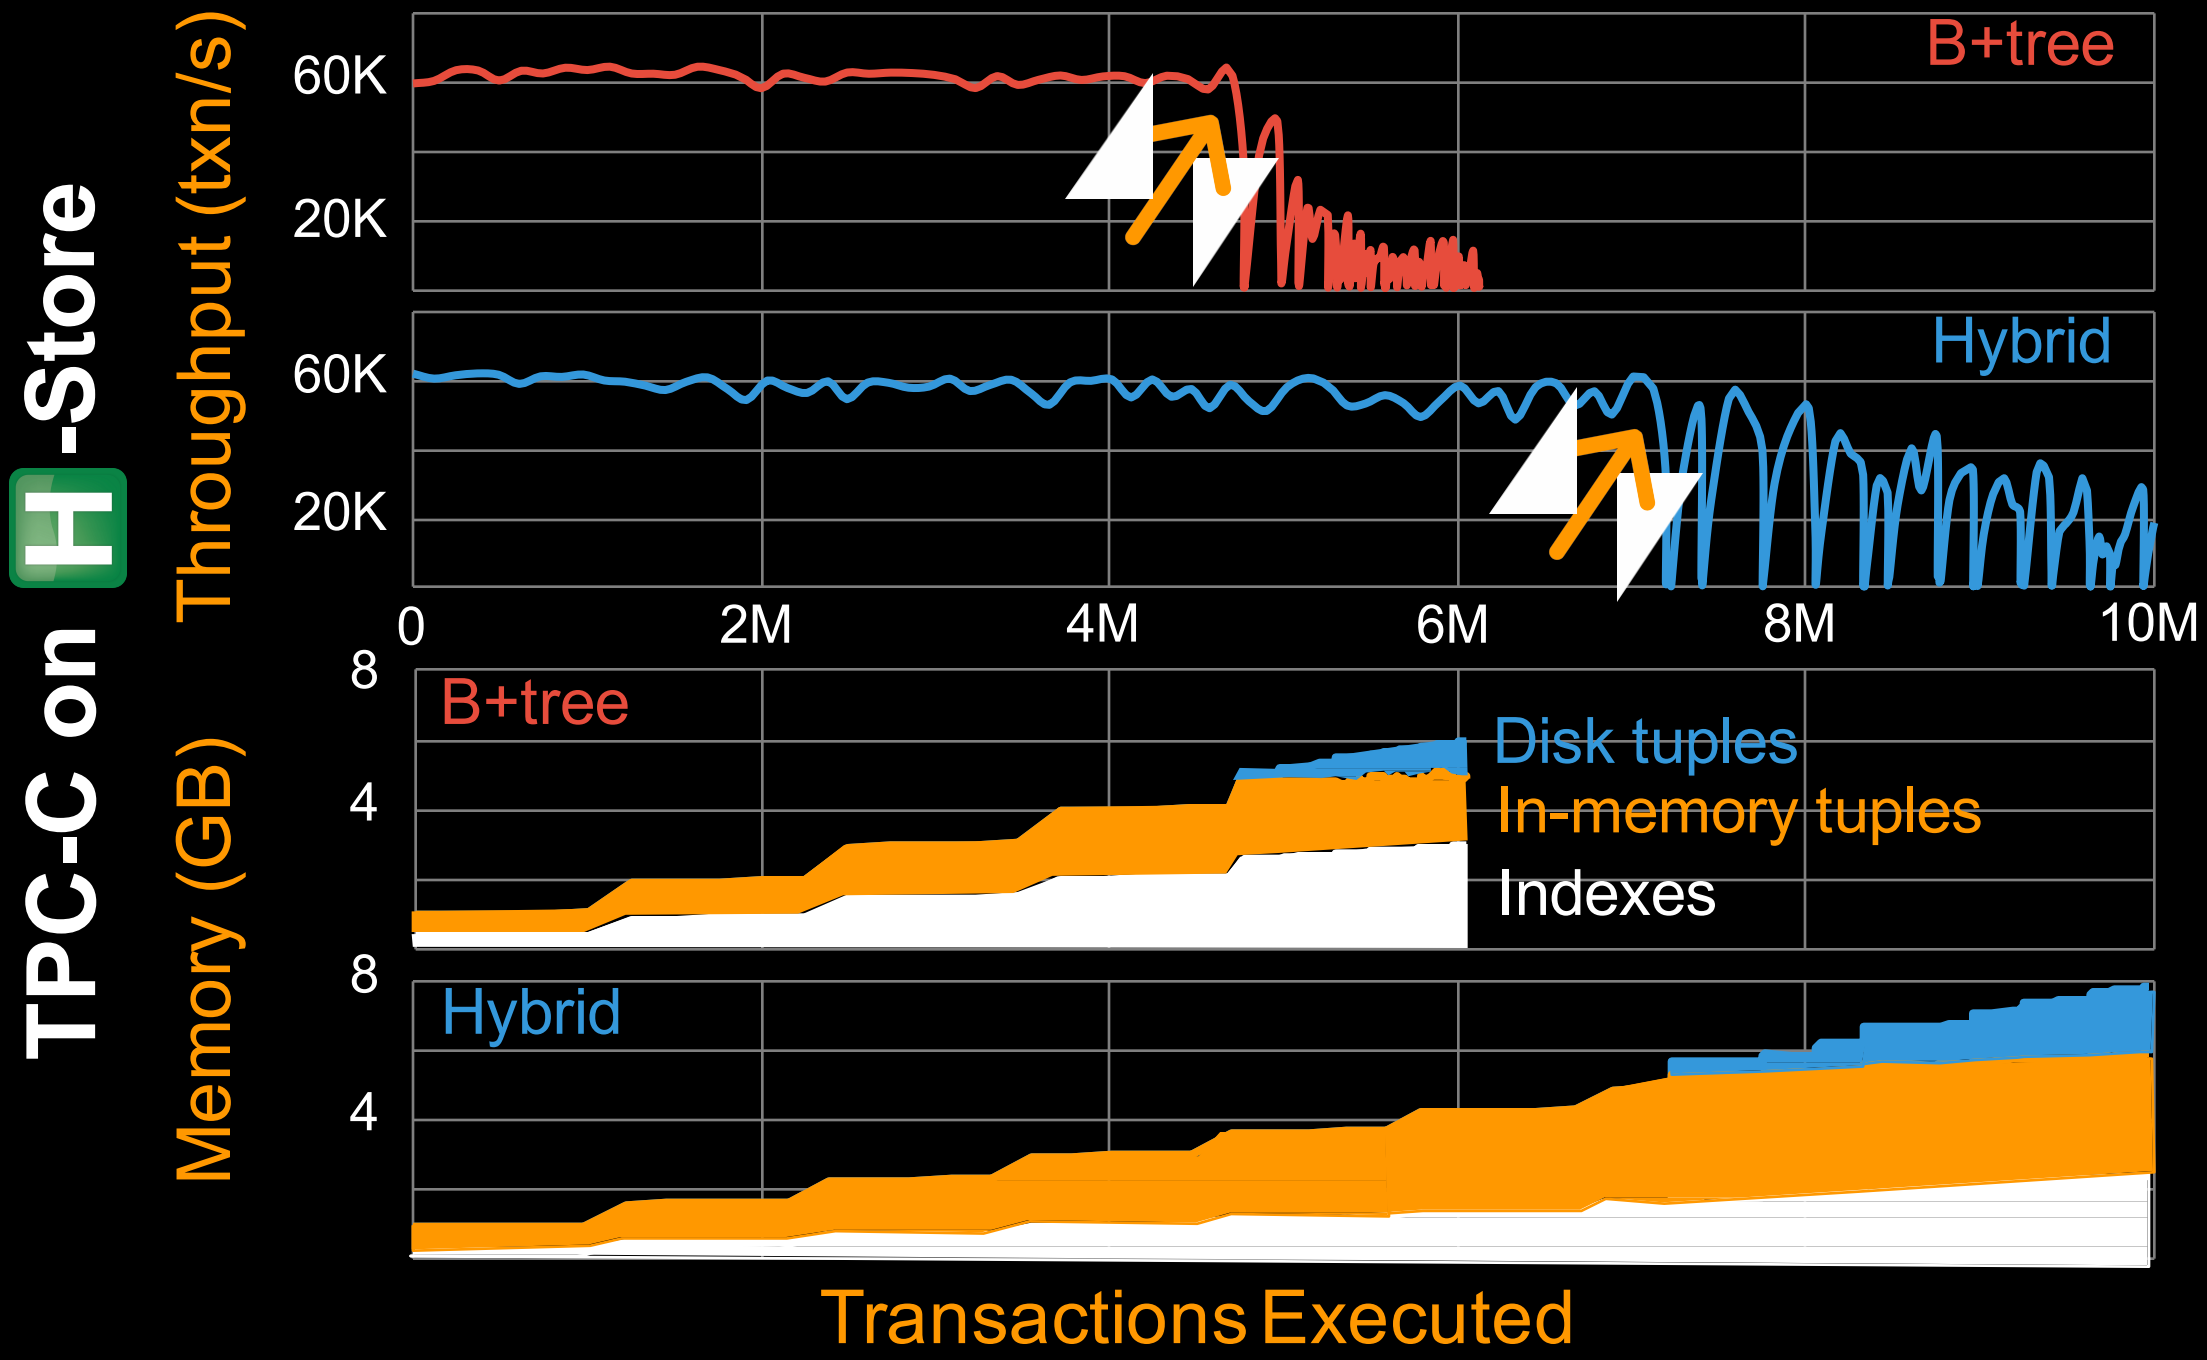
TPC-C on H-Store
Throughput (txn/s)
60K
20K
B+tree
60K
20K
Hybrid
0
2M
4M
6M
8M
10M
Memory (GB)
8
4
B+tree
Disk tuples
In-memory tuples
Indexes
8
4
Hybrid
Transactions Executed

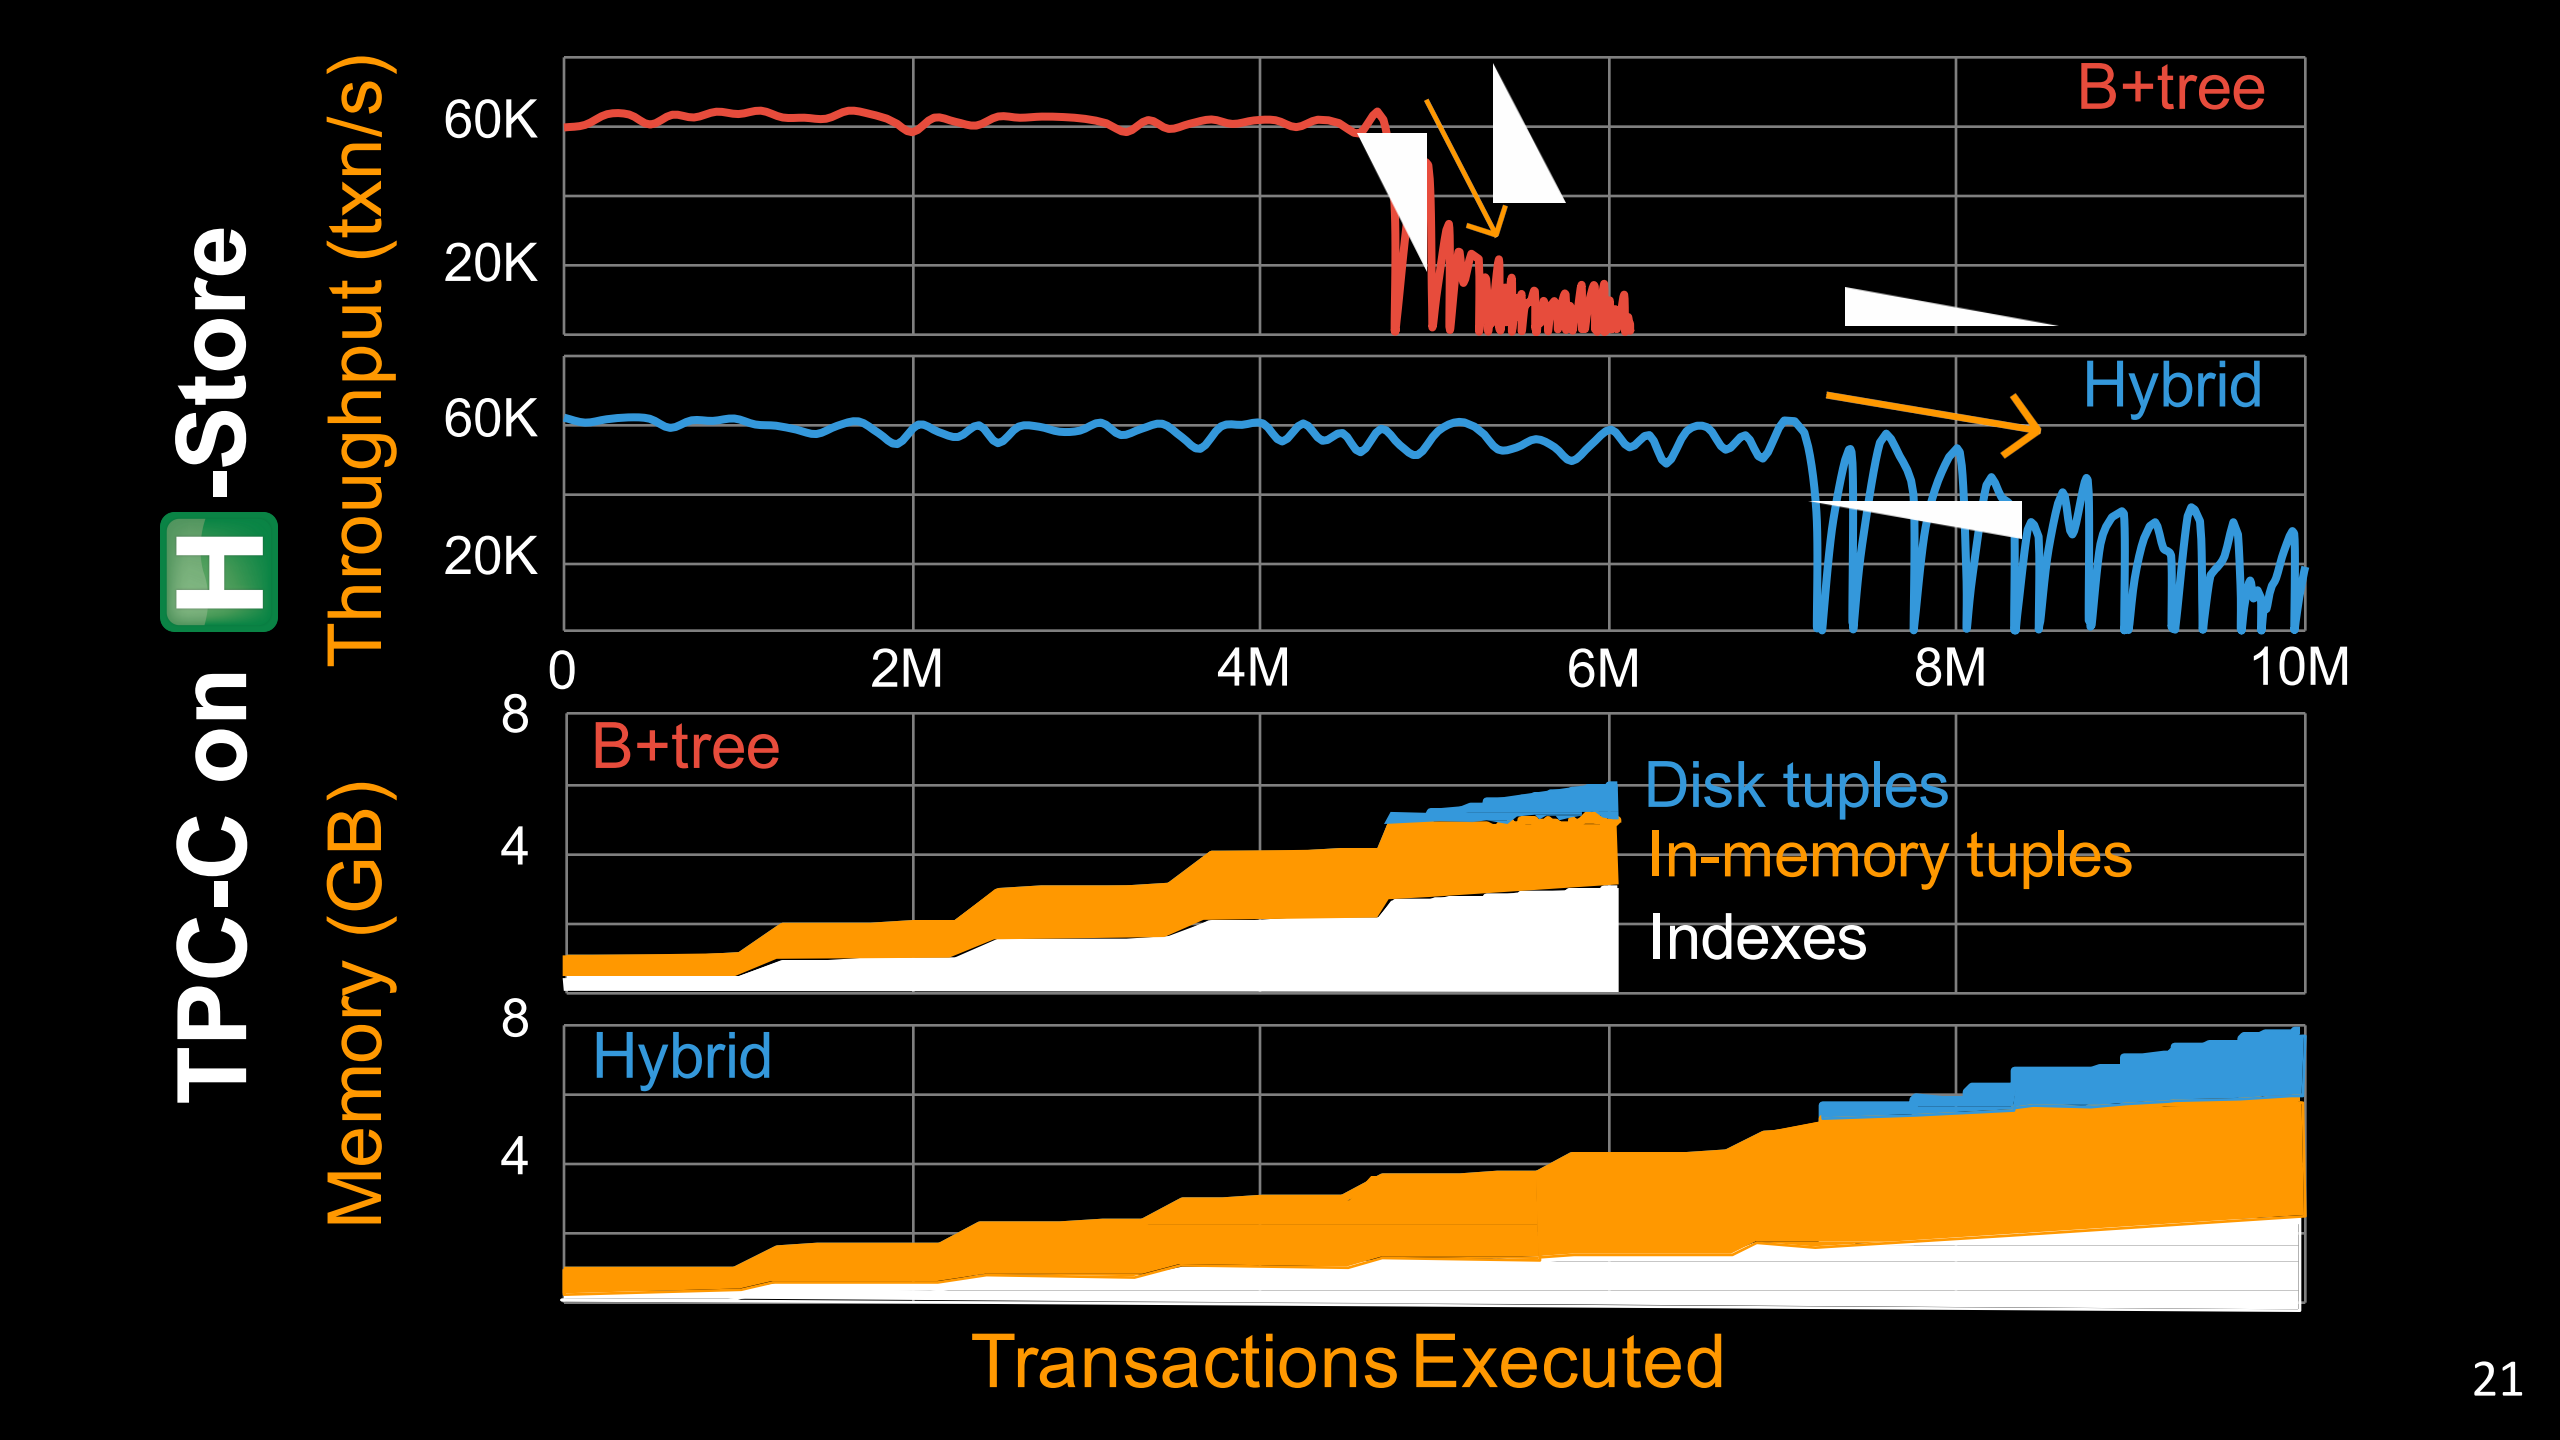
TPC-C on H-Store
Throughput (txn/s)
60K
20K
B+tree
60K
20K
Hybrid
0
2M
4M
6M
8M
10M
Memory (GB)
8
4
B+tree
Disk tuples
In-memory tuples
Indexes
8
4
Hybrid
Transactions Executed

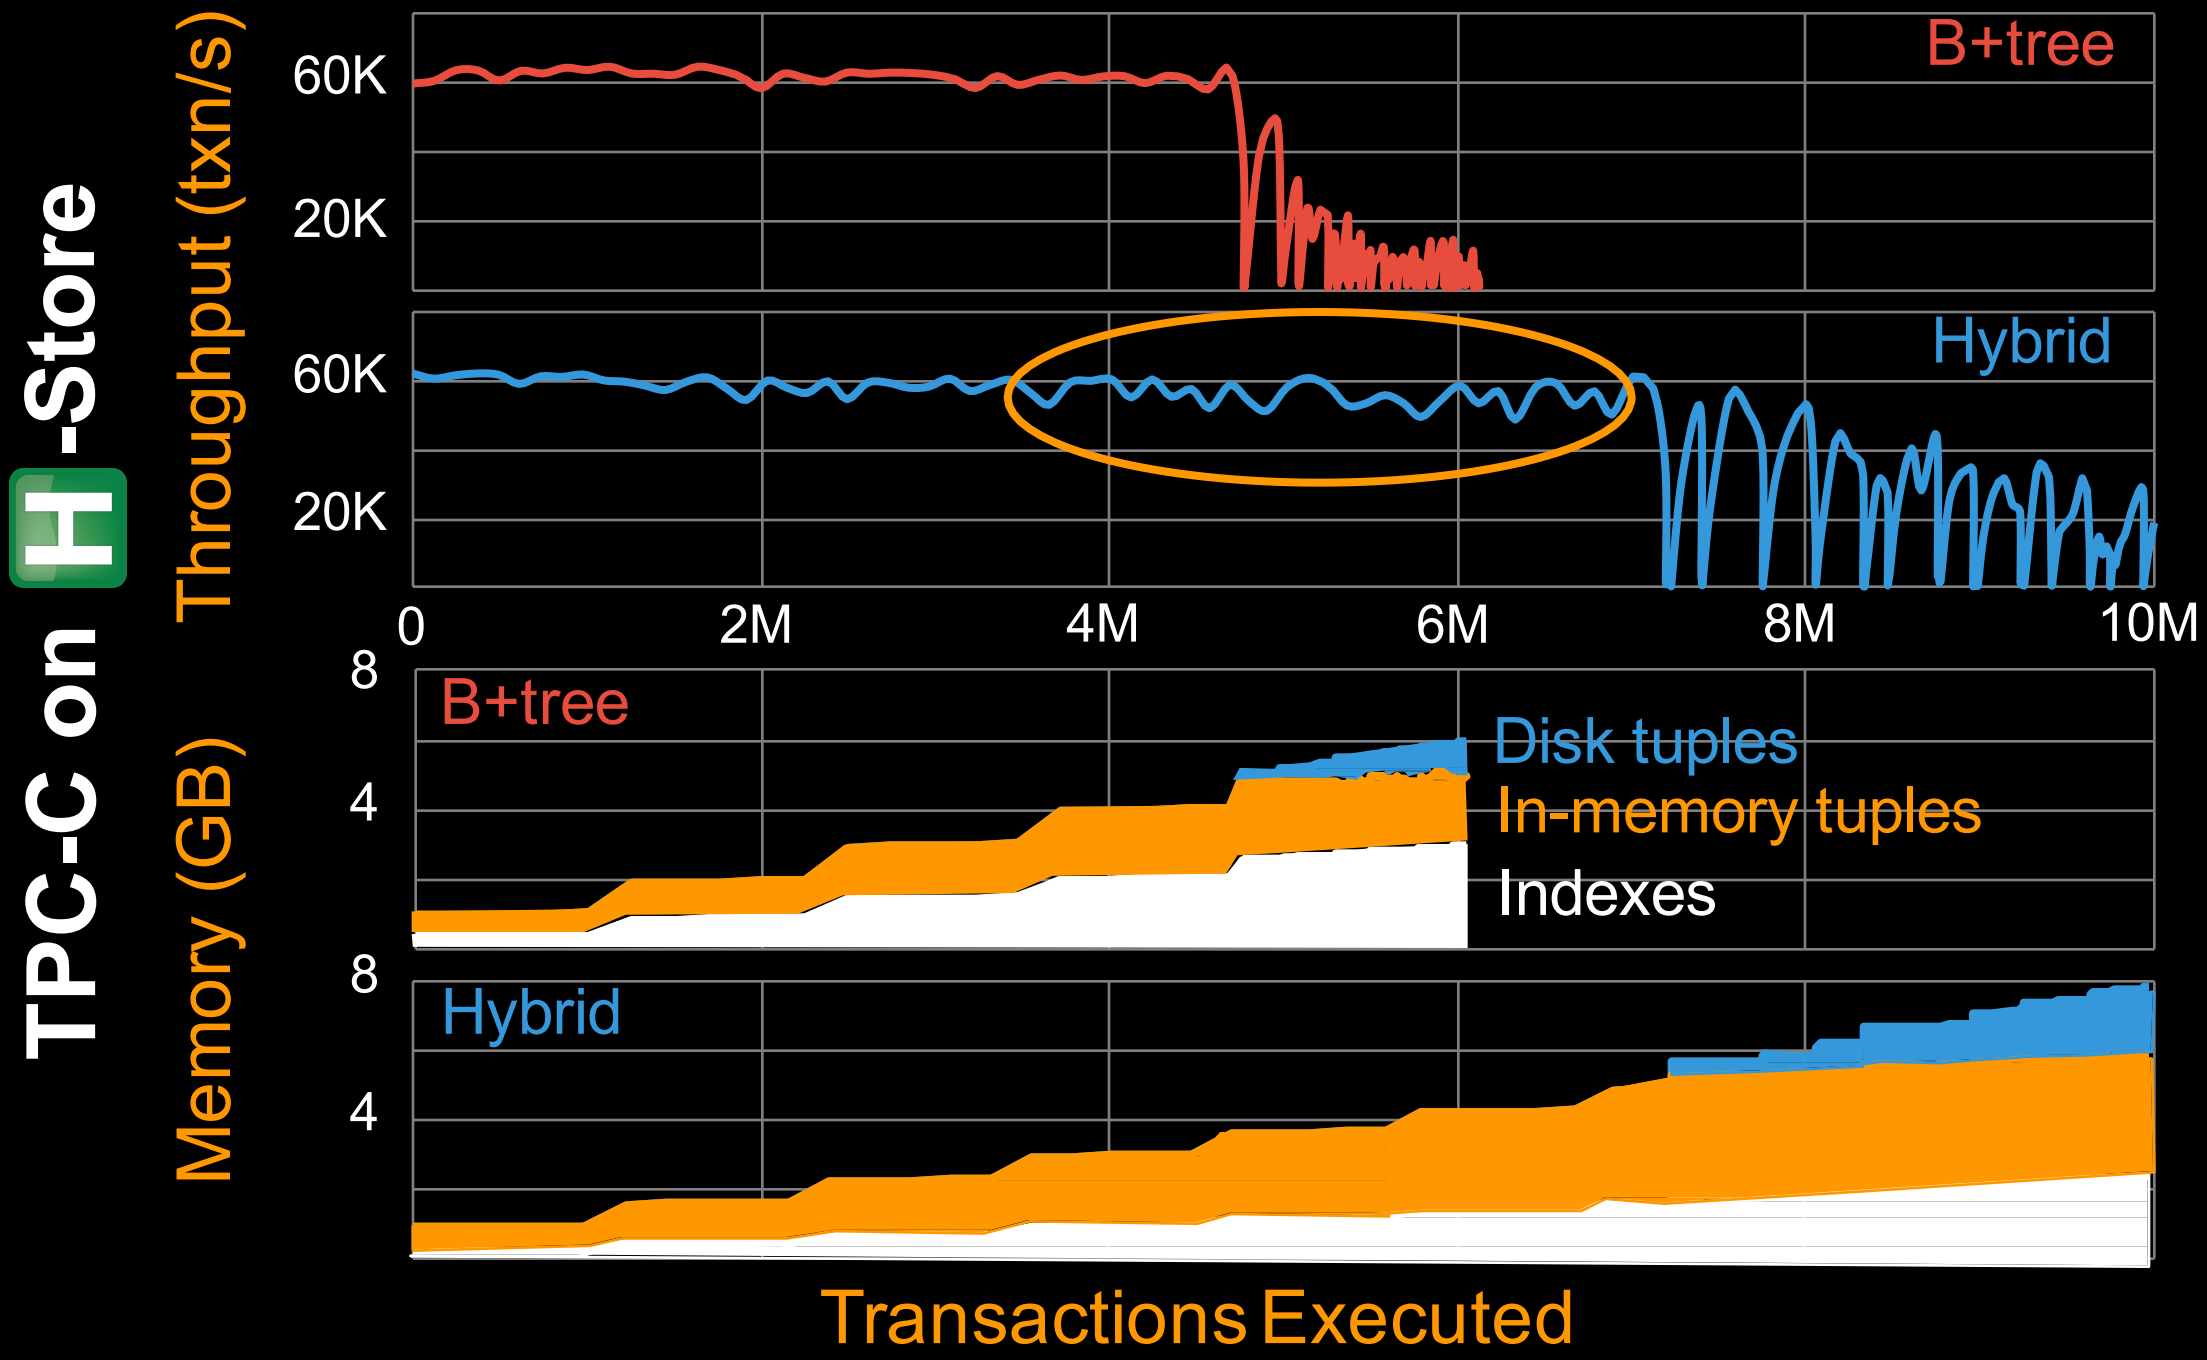
TPC-C on H-Store
Throughput (txn/s)
60K
20K
B+tree
60K
20K
Hybrid
0
2M
4M
6M
8M
10M
Memory (GB)
8
4
B+tree
Disk tuples
In-memory tuples
Indexes
8
4
Hybrid
Transactions Executed

Take Away:

Memory saved by indexes → **Larger working set in memory** → **Higher throughput**

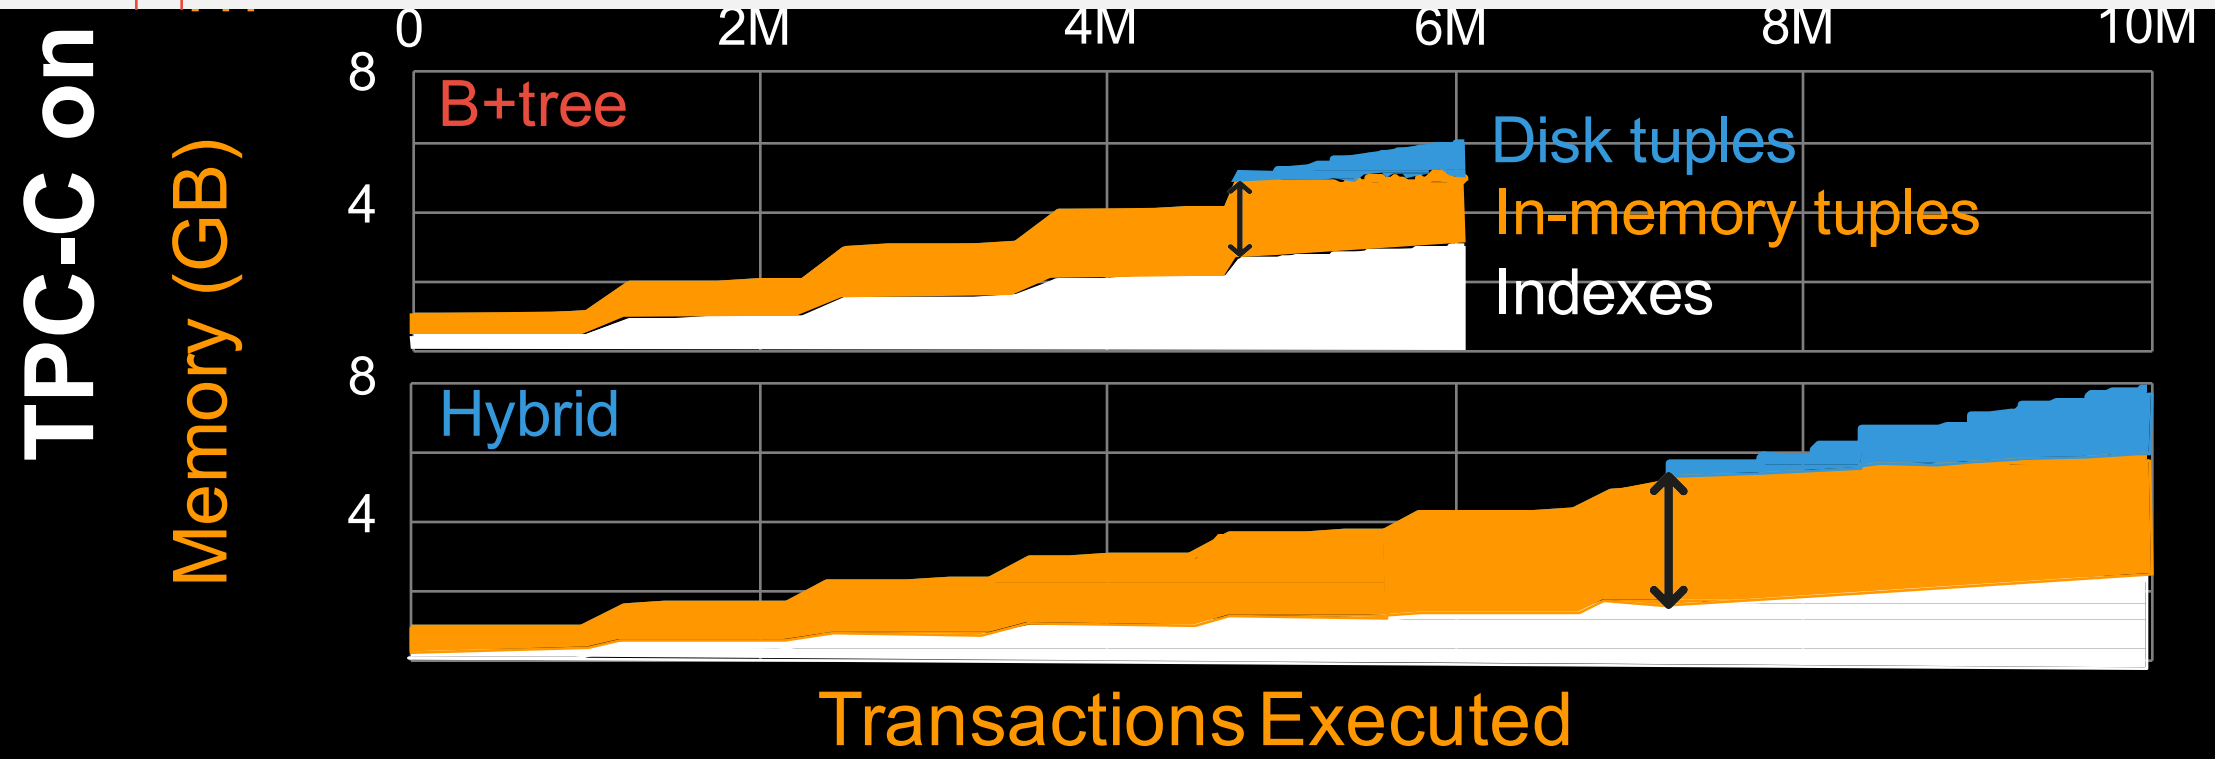

This is just the BEGINNING

Conclusions

① The hybrid index architecture **GENERAL**

② The Dual-Stage Transformation **PRACTICAL**

③ Applied to 4 index structures **USEFUL**

- B+tree
- Masstree
- Skip List
- Adaptive Radix Tree (ART)

Hybrid Index Inspirations:

1. A Tradeoff Research among Data Tuples, Indexes and Evicted Tables in Memory Consumption
2. Dual-stage Architecture of Indexes in Hybrid Memory
3. A Memory-efficient Hash Table Index with Range Query Optimization (DHT might be possible)
4. Non-blocking Merging for Hybrid Indexes with COW

Anti-caching Inspirations:

1. Revisit Anti-caching Mechanism for OLTP Workloads in Hybrid Memory (Evict cold tuples to NVM)
2. Non-blocking Eviction in NVM-Optimized Anti-caching

Non-volatile Memory Inspirations:

1. DiRedis: A NVM-Optimized KV-Store Based on Redis
2. Rethinking Program Scheme in Persistent Memory Era

Thanks !

Note website:
http://kaixinhuang.com/Research/hybrid-index-db/

The Art of Research Presentation